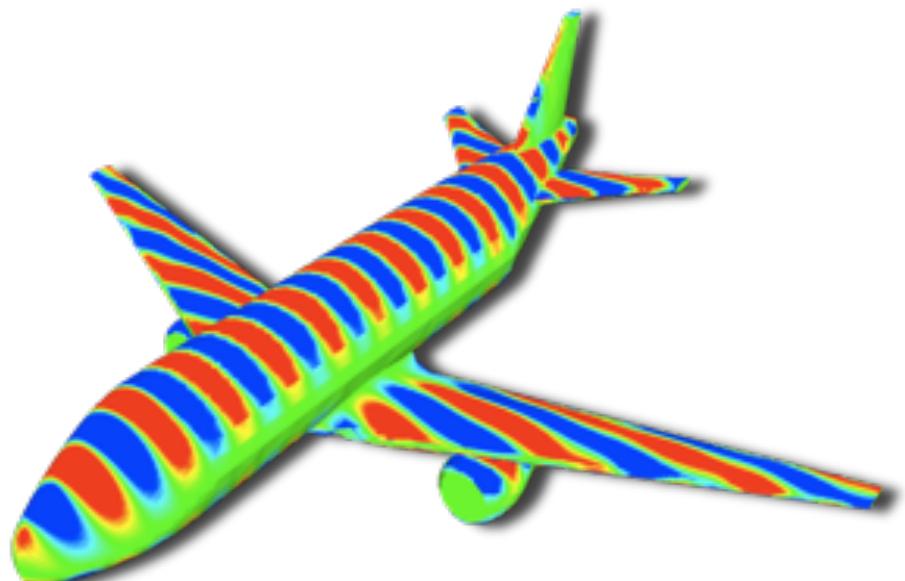
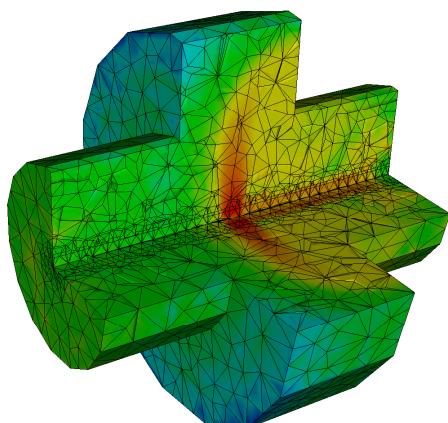




# DG-FEM for PDE's

## Lecture 8

Jan S Hesthaven  
Brown University  
[Jan.Hesthaven@Brown.edu](mailto:Jan.Hesthaven@Brown.edu)



# A brief overview of what's to come

---

- ▶ Lecture 1: Introduction and DG in 1D
- ▶ Lecture 2: Implementation and numerical aspects
- ▶ Lecture 3: Insight through theory
- ▶ Lecture 4: Nonlinear problems
- ▶ Lecture 5/6: 2D DG-FEM, implementation and applications
- ▶ Lecture 7: Higher order/Global problems
- ▶ **Lecture 8: 3D and advanced topics**

# Lecture 8

---

- ✓ Let's briefly recall what we know
- ✓ Part I: 3D problems and extensions
  - ✓ Formulations and examples
  - ✓ Adaptivity and curvilinear elements
- ✓ Part II: The need for speed
  - ✓ Parallel computing
  - ✓ GPU computing
  - ✓ Software beyond Matlab

# Lets summarize

---

We are done with all the basics -- and we have started to see it work for us -- we know how to do

- ✓ 1D/2D problems
- ✓ Linear/nonlinear problems
- ✓ First and higher operators
- ✓ Complex geometries
- ✓ ... and we have insight into theory

All we need is 3D -- and with that comes  
the need for speed !

# Extension to 3D ?

It is really simple at this stage !

Weak form:

$$\int_{\mathbf{D}^k} \left[ \frac{\partial u_h^k}{\partial t} \ell_n^k(\mathbf{x}) - \mathbf{f}_h^k \cdot \nabla \ell_n^k(\mathbf{x}) \right] d\mathbf{x} = - \oint_{\partial \mathbf{D}^k} \hat{\mathbf{n}} \cdot \mathbf{f}^* \ell_n^k(\mathbf{x}) d\mathbf{x},$$

Strong form:

$$\int_{\mathbf{D}^k} \left[ \frac{\partial u_h^k}{\partial t} + \nabla \cdot \mathbf{f}_h^k \right] \ell_n^k(\mathbf{x}) d\mathbf{x} = \oint_{\partial \mathbf{D}^k} \hat{\mathbf{n}} \cdot \left[ \mathbf{f}_h^k - \mathbf{f}^* \right] \ell_n^k(\mathbf{x}) d\mathbf{x},$$

$$\mathbf{f}^* = \{\{\mathbf{f}_h(\mathbf{u}_h)\}\} + \frac{C}{2} [\![\mathbf{u}_h]\!]. \quad C = \max_u \left| \lambda \left( \hat{\mathbf{n}} \cdot \frac{\partial \mathbf{f}}{\partial \mathbf{u}} \right) \right|,$$

**Nothing is essential new**

# Extension to 3D

---

Apart from the ‘logistics’ all we need to worry about is to choose our element and how to represent the solution

$$u(\mathbf{r}) \simeq u_h(\mathbf{r}) = \sum_{n=1}^{N_p} \hat{u}_n \psi_n(\mathbf{r}) = \sum_{i=1}^{N_p} u(\mathbf{r}_i) \ell_i(\mathbf{r}),$$
$$\mathbf{u} = \mathcal{V} \hat{\mathbf{u}}, \quad \mathcal{V}^T \ell(\mathbf{r}) = \boldsymbol{\psi}(\mathbf{r}), \quad \mathcal{V}_{ij} = \psi_j(\mathbf{r}_i).$$

We need points

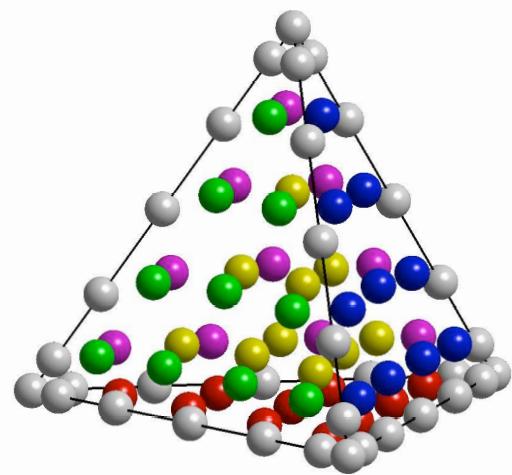
$$N_p = \frac{(N+1)(N+2)(N+3)}{6},$$

We need an orthonormal basis

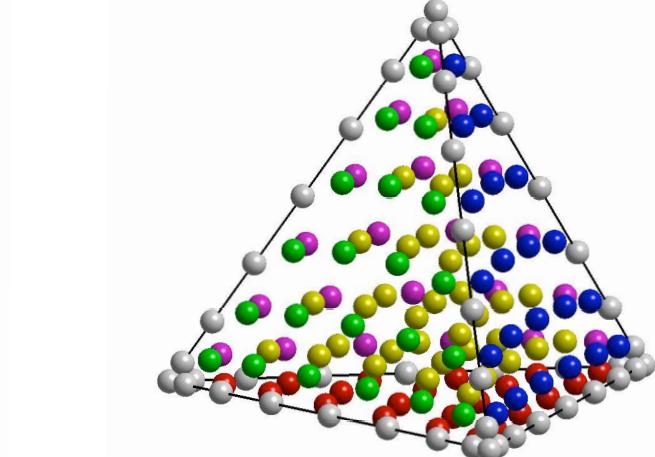
$$\psi_{ijk}(r, s, t) = 2\sqrt{2} P_i^{(0,0)}(a) P_j^{(2i+1,0)}(b) P_k^{(2i+2j+2,0)}(b) (1-b)^i (1-c)^{i+j},$$

# Extension to 3D

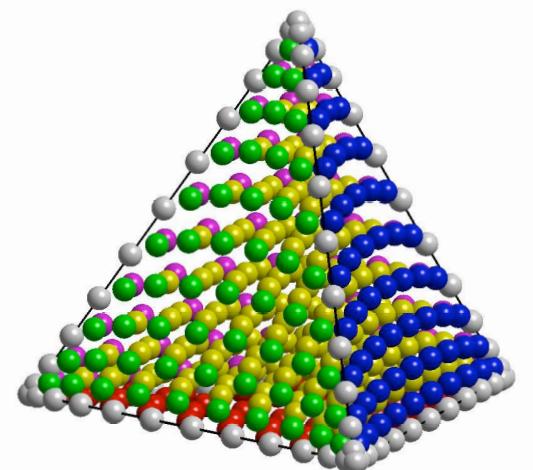
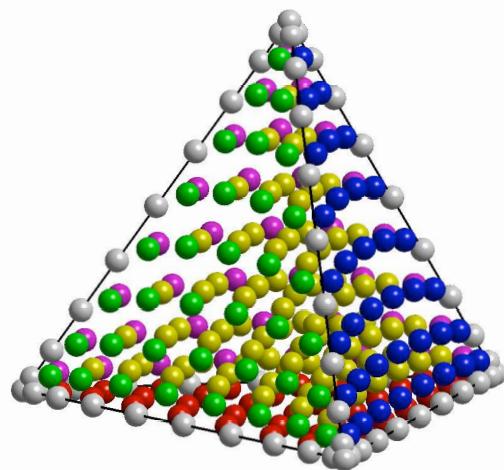
For other element types, one simply need to define nodes and modes for that elements



: 6



$N = 8$



# Extension to 3D

---

Everything is identical in spirit

Mass matrix

$$\mathcal{M}^k = J^k (\mathcal{V} \mathcal{V}^T)^{-1}.$$

Diff matrix

$$\mathcal{D}_r \mathcal{V} = \mathcal{V}_r, \quad \mathcal{D}_s \mathcal{V} = \mathcal{V}_s, \quad \mathcal{D}_t \mathcal{V} = \mathcal{V}_t,$$

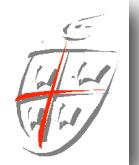
Derivative

$$\begin{aligned}\frac{\partial}{\partial x} &= \frac{\partial r}{\partial x} \mathcal{D}_r + \frac{\partial s}{\partial x} \mathcal{D}_s + \frac{\partial t}{\partial x} \mathcal{D}_t, \\ \frac{\partial}{\partial y} &= \frac{\partial r}{\partial y} \mathcal{D}_r + \frac{\partial s}{\partial y} \mathcal{D}_s + \frac{\partial t}{\partial y} \mathcal{D}_t, \\ \frac{\partial}{\partial z} &= \frac{\partial r}{\partial z} \mathcal{D}_r + \frac{\partial s}{\partial z} \mathcal{D}_s + \frac{\partial t}{\partial z} \mathcal{D}_t,\end{aligned}$$

Stiffness matrix

$$\mathcal{S}_r = \mathcal{M}^{-1} \mathcal{D}_r, \quad \mathcal{S}_s = \mathcal{M}^{-1} \mathcal{D}_s, \quad \mathcal{S}_t = \mathcal{M}^{-1} \mathcal{D}_t.$$

# Example - Maxwell's equations



Consider Maxwell's equations

$$\varepsilon \partial_t E - \nabla \times H = -j, \quad \mu \partial_t H + \nabla \times E = 0,$$

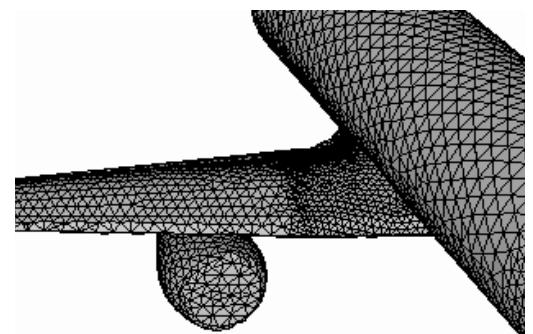
Write it on conservation form as

$$\frac{\partial q}{\partial t} + \nabla \cdot F = -J \quad F = \begin{bmatrix} -\hat{e} \times H \\ \hat{e} \times E \end{bmatrix} \quad q = \begin{bmatrix} E \\ H \end{bmatrix}$$

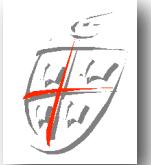
Represent the solution as

$$\Omega = \sum_k D^k \quad q_N = \sum_{i=1}^N q(\mathbf{x}_i, t) L_i(\mathbf{x})$$

and assume



$$\int_D \left( \frac{\partial \mathbf{q}_N}{\partial t} + \nabla \cdot \mathbf{F}_N - \mathbf{J}_N \right) L_i(\mathbf{x}) \, d\mathbf{x} = \oint_{\partial D} L_i(\mathbf{x}) \hat{\mathbf{n}} \cdot [\mathbf{F}_N - \mathbf{F}^*] \, d\mathbf{x}.$$



# Example - Maxwell's equations

On each element we then define

$$\hat{M}_{ij} = \int_D L_i L_j \, dx, \quad \hat{S}_{ij} = \int_D \nabla L_j L_i \, dx, \quad \hat{F}_{ij} = \oint_{\partial D} L_i L_j \, dx,$$

With the numerical flux given as

$$\hat{\mathbf{n}} \cdot [\mathbf{F} - \mathbf{F}^*] = \begin{cases} \mathbf{n} \times (\gamma \mathbf{n} \times [\mathbf{E}] - [\mathbf{B}]), \\ \mathbf{n} \times (\gamma \mathbf{n} \times [\mathbf{B}] + [\mathbf{E}]), \end{cases} \quad [Q] = Q^- - Q^+$$

To obtain the local matrix based scheme

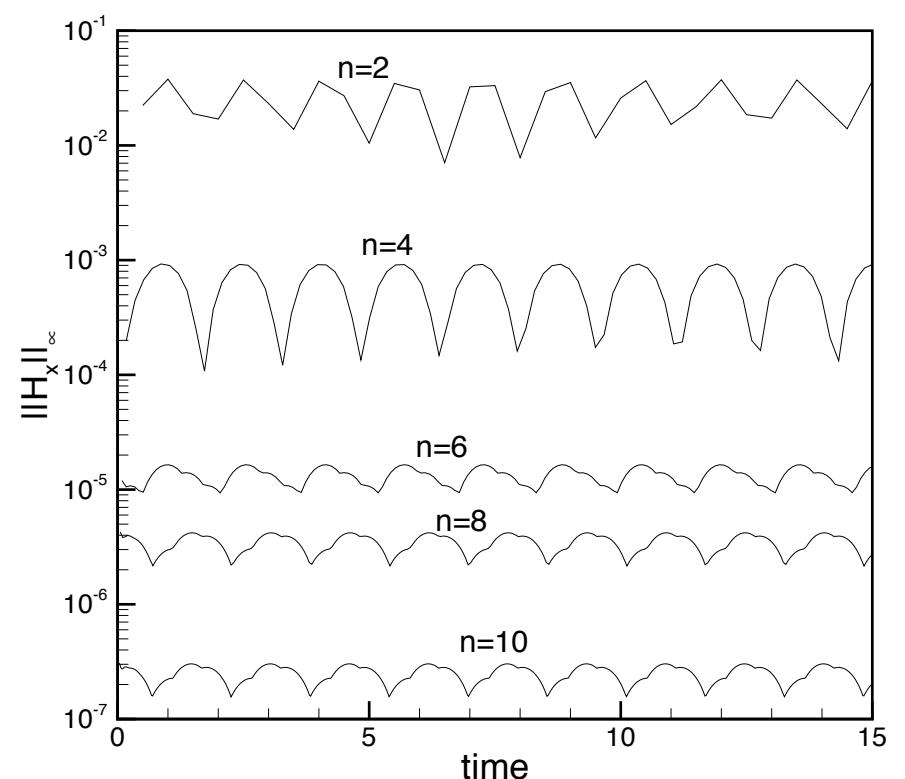
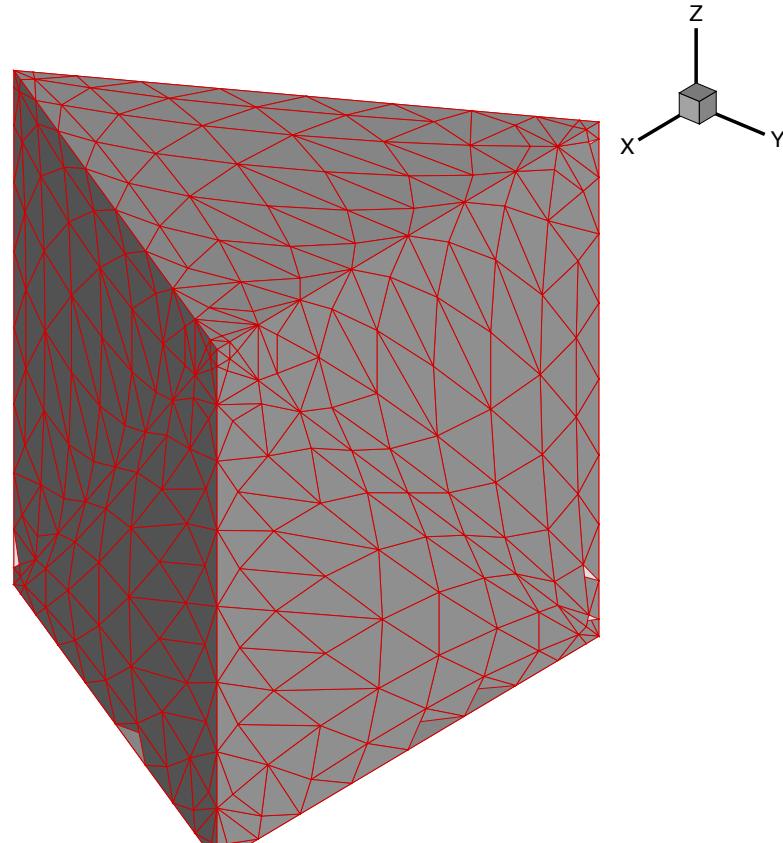
$$\hat{M} \frac{d\hat{\mathbf{q}}}{dt} + \hat{S} \cdot \hat{\mathbf{F}} - \hat{M} \hat{\mathbf{J}} = \hat{\mathbf{F}} \hat{\mathbf{n}} \cdot [\hat{\mathbf{F}} - \hat{\mathbf{F}}^*],$$

One then typically uses an explicit Runge-Kutta to advance in time - just like 1D/2D.

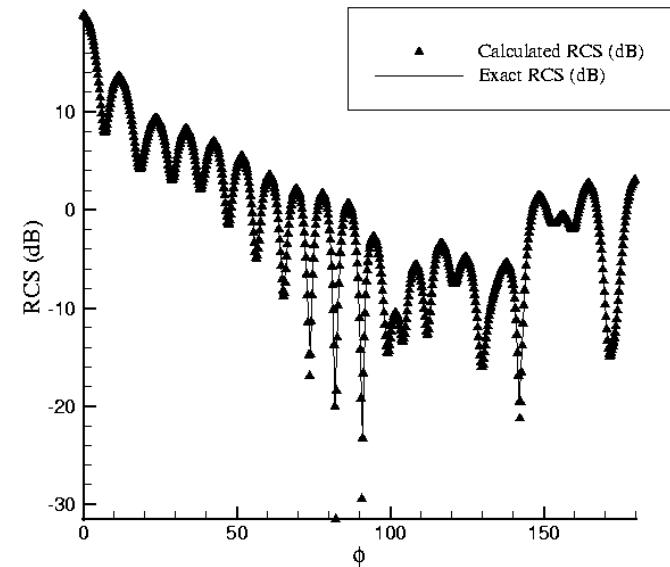
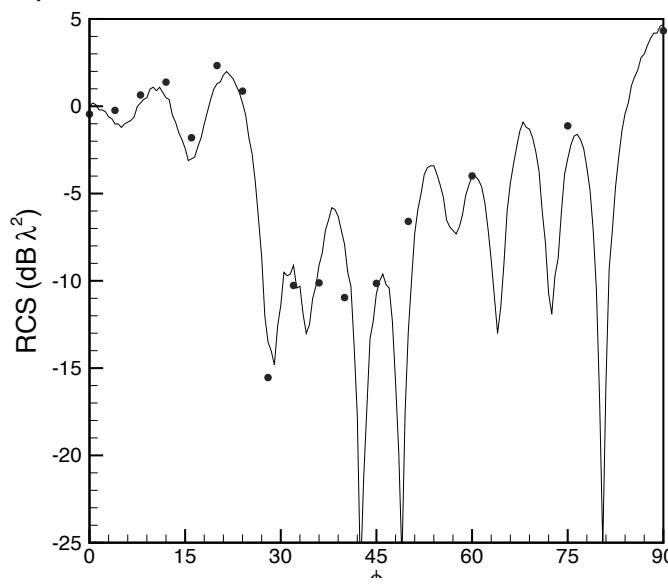
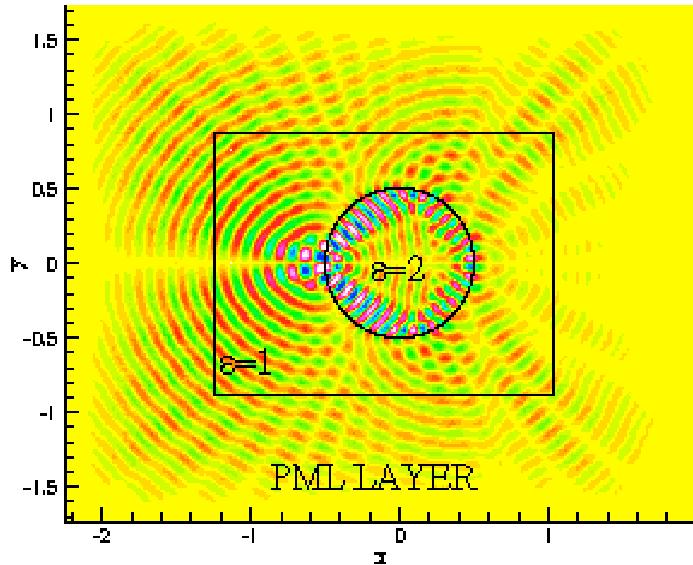
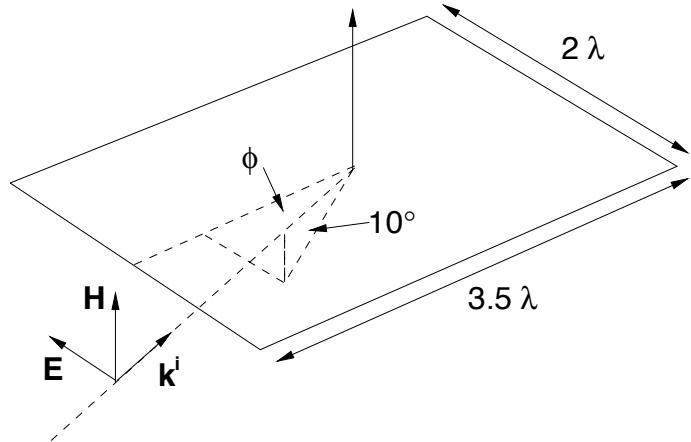
# An example - Maxwell's equations

---

## Simple wave propagation

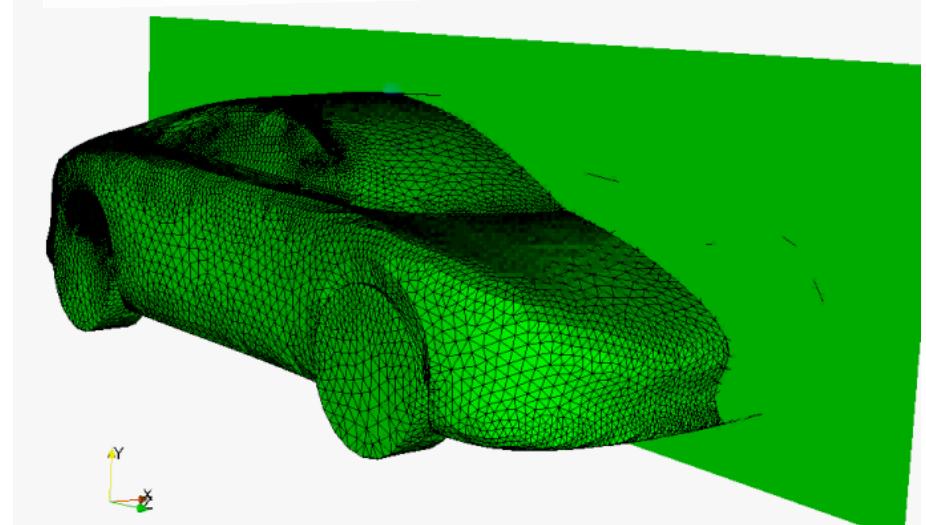
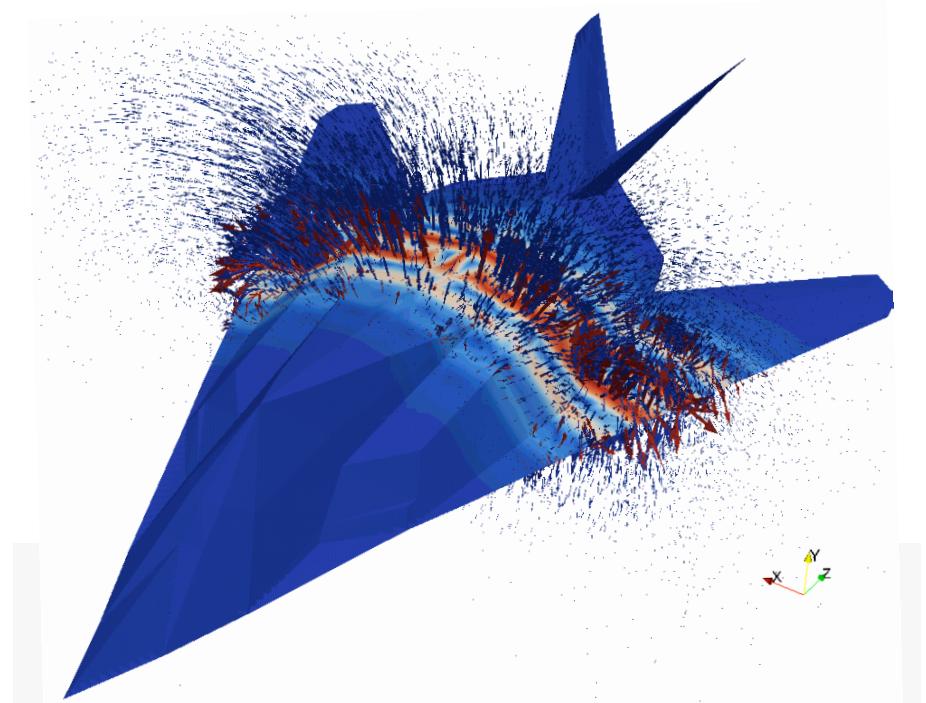
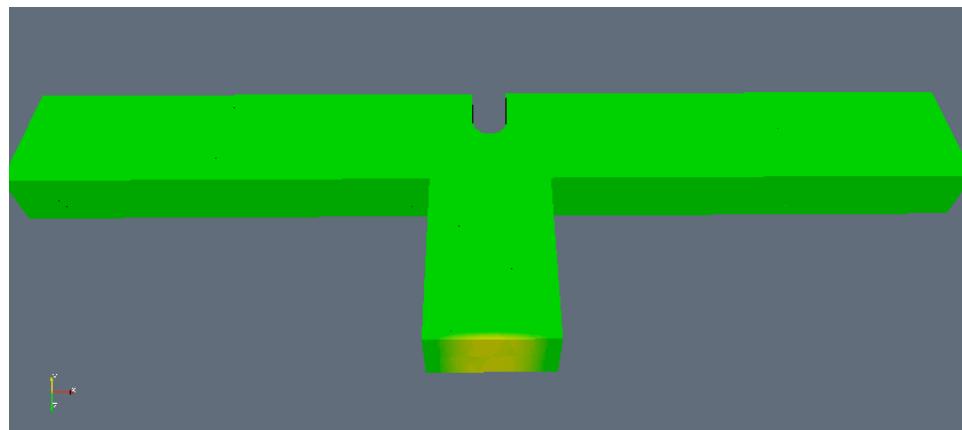
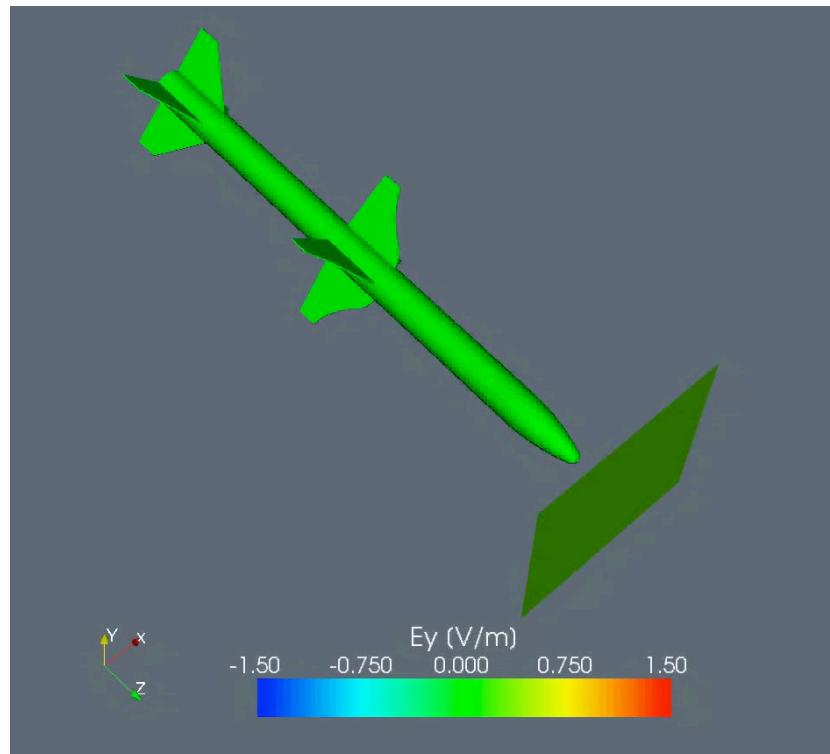


# An example - Maxwell's equations



# An example - Maxwell's equations

---



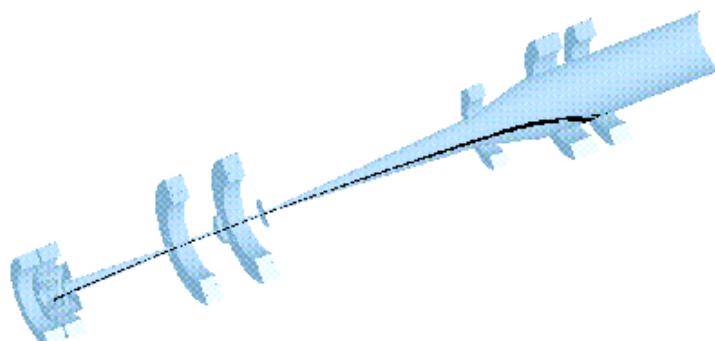
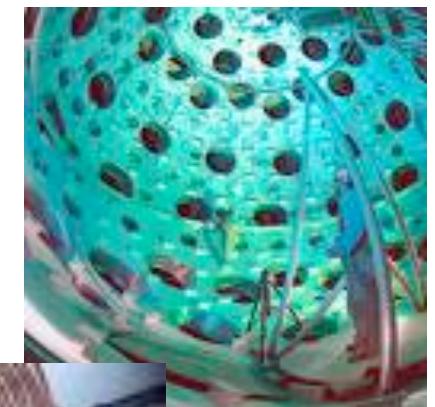
Animations by Nico Godel (Hamburg)

# Kinetic Plasma Physics

---

## Important applications

- ✓ High-power/High-frequency microwave generation
- ✓ Particle accelerators
- ✓ Laser-matter interaction
- ✓ Fusion applications, e.g., plasma edge
- ✓ etc



# Kinetic Plasma Physics

In high-speed plasma problems dominated by kinetic effects, one needs to solve for  $f(x,p,t)$  - 6D+1

*Vlasov/Boltzmann equation*

$$\partial_t f + \mathbf{v} \cdot \partial_x f + q(\mathbf{E} + \mathbf{v} \times \mathbf{B}) \cdot \partial_p f = \langle \text{Sources} \rangle - \langle \text{Sinks} \rangle.$$

*Maxwell's equations*

$$\partial_t \mathbf{E} - \frac{1}{\epsilon} \nabla \times \mathbf{H} = -\frac{\mathbf{j}}{\epsilon},$$

$$\partial_t \mathbf{H} + \frac{1}{\mu} \nabla \times \mathbf{E} = 0,$$

$$\nabla \cdot \mathbf{H} = 0, \quad \nabla \cdot \mathbf{E} = \frac{\rho}{\epsilon}.$$

Coupled through

$$\rho := \int f \, dv, \quad j := \int v f \, dv.$$

# Particle-in-Cell (PIC) Methods

---

This is an attempt to solve the Vlasov/Boltzmann equation by sampling with P particles

$$f(x, p, t) = \sum_{n=1}^P q_n S(x - x_n(t)) \delta(p - p_n(t)),$$

$$\rho(x, t) = \sum_{n=1}^P q_n S(x - x_n(t)), \quad j(x, t) = \sum_{n=1}^P v_n q_n S(x - x_n(t))$$

Ideally we have

$$S(x) = \delta(x) \quad \leftarrow \quad \text{a point particle}$$

However, this is not practical, nor reasonable - so  $S(x)$  is a **shape-function**

# Particle-in-Cell Methods

## Maxwell's equations

$$\begin{aligned}\varepsilon \partial_t E - \nabla \times H &= -j, & \mu \partial_t H + \nabla \times E &= 0, \\ \nabla \cdot (\varepsilon E) &= \rho, & \nabla \cdot (\mu H) &= 0,\end{aligned}$$

## Particle/Phase dynamics

$$\frac{dx_n}{dt} = v_n(t) \quad \frac{dmv_n}{dt} = q_n(E + v_n \times H) \quad m = \frac{1}{\sqrt{1 - (v_n/c)^2}}$$

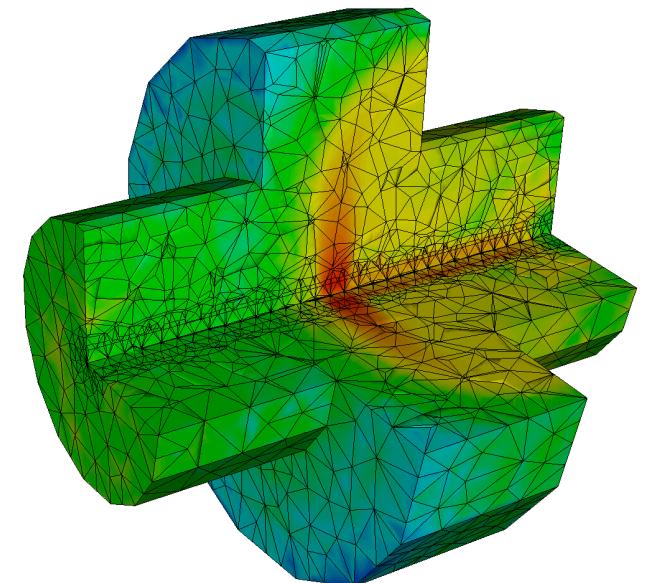
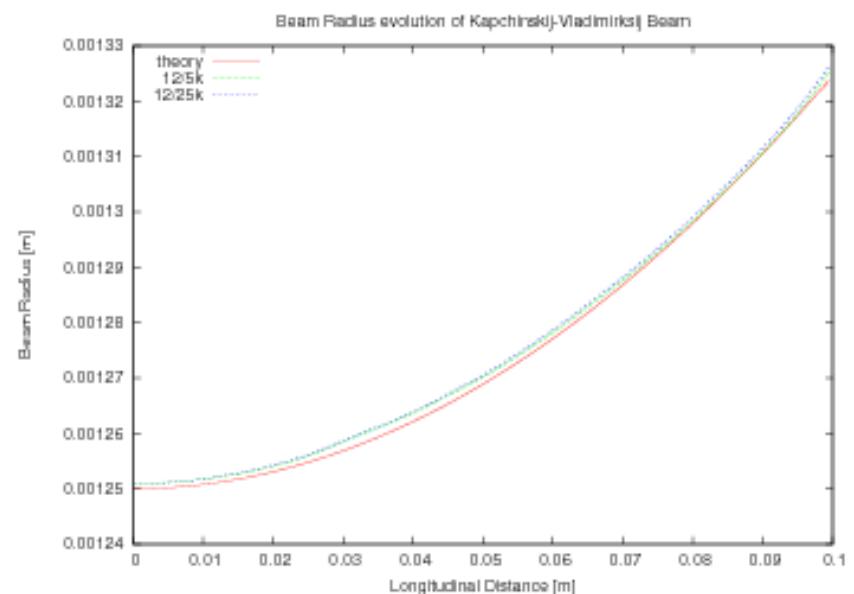
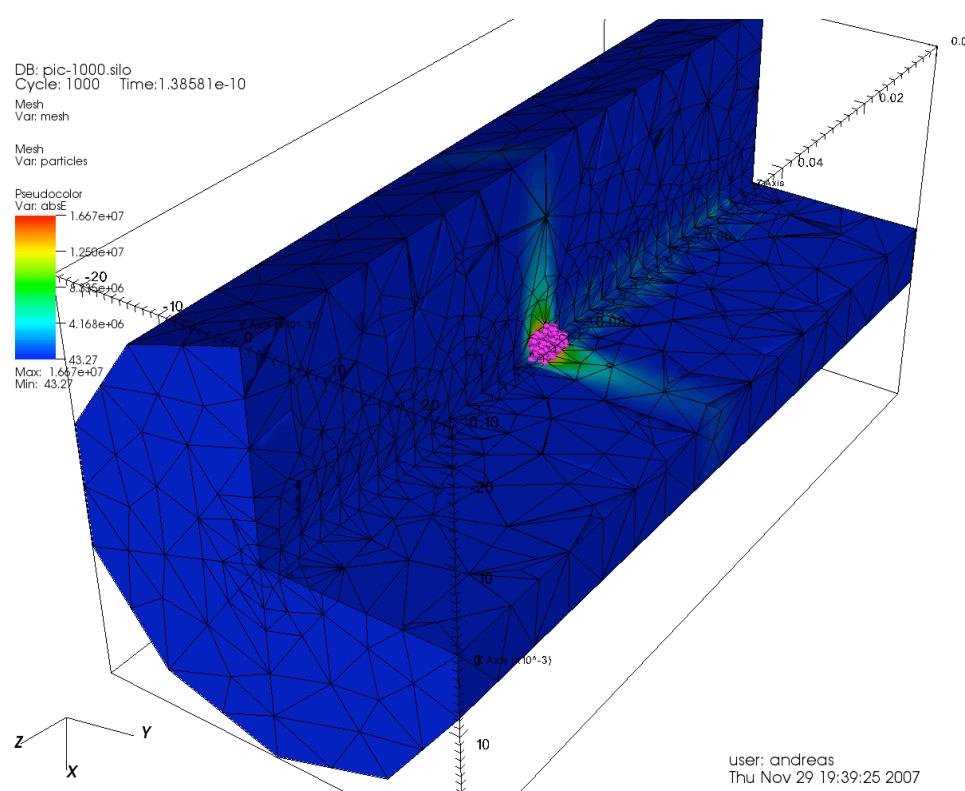
## Particles-to-fields

$$\rho(x, t) = \sum_{n=1}^P q_n S(x - x_n(t)), \quad j(x, t) = \sum_{n=1}^P v_n q_n S(x - x_n(t))$$

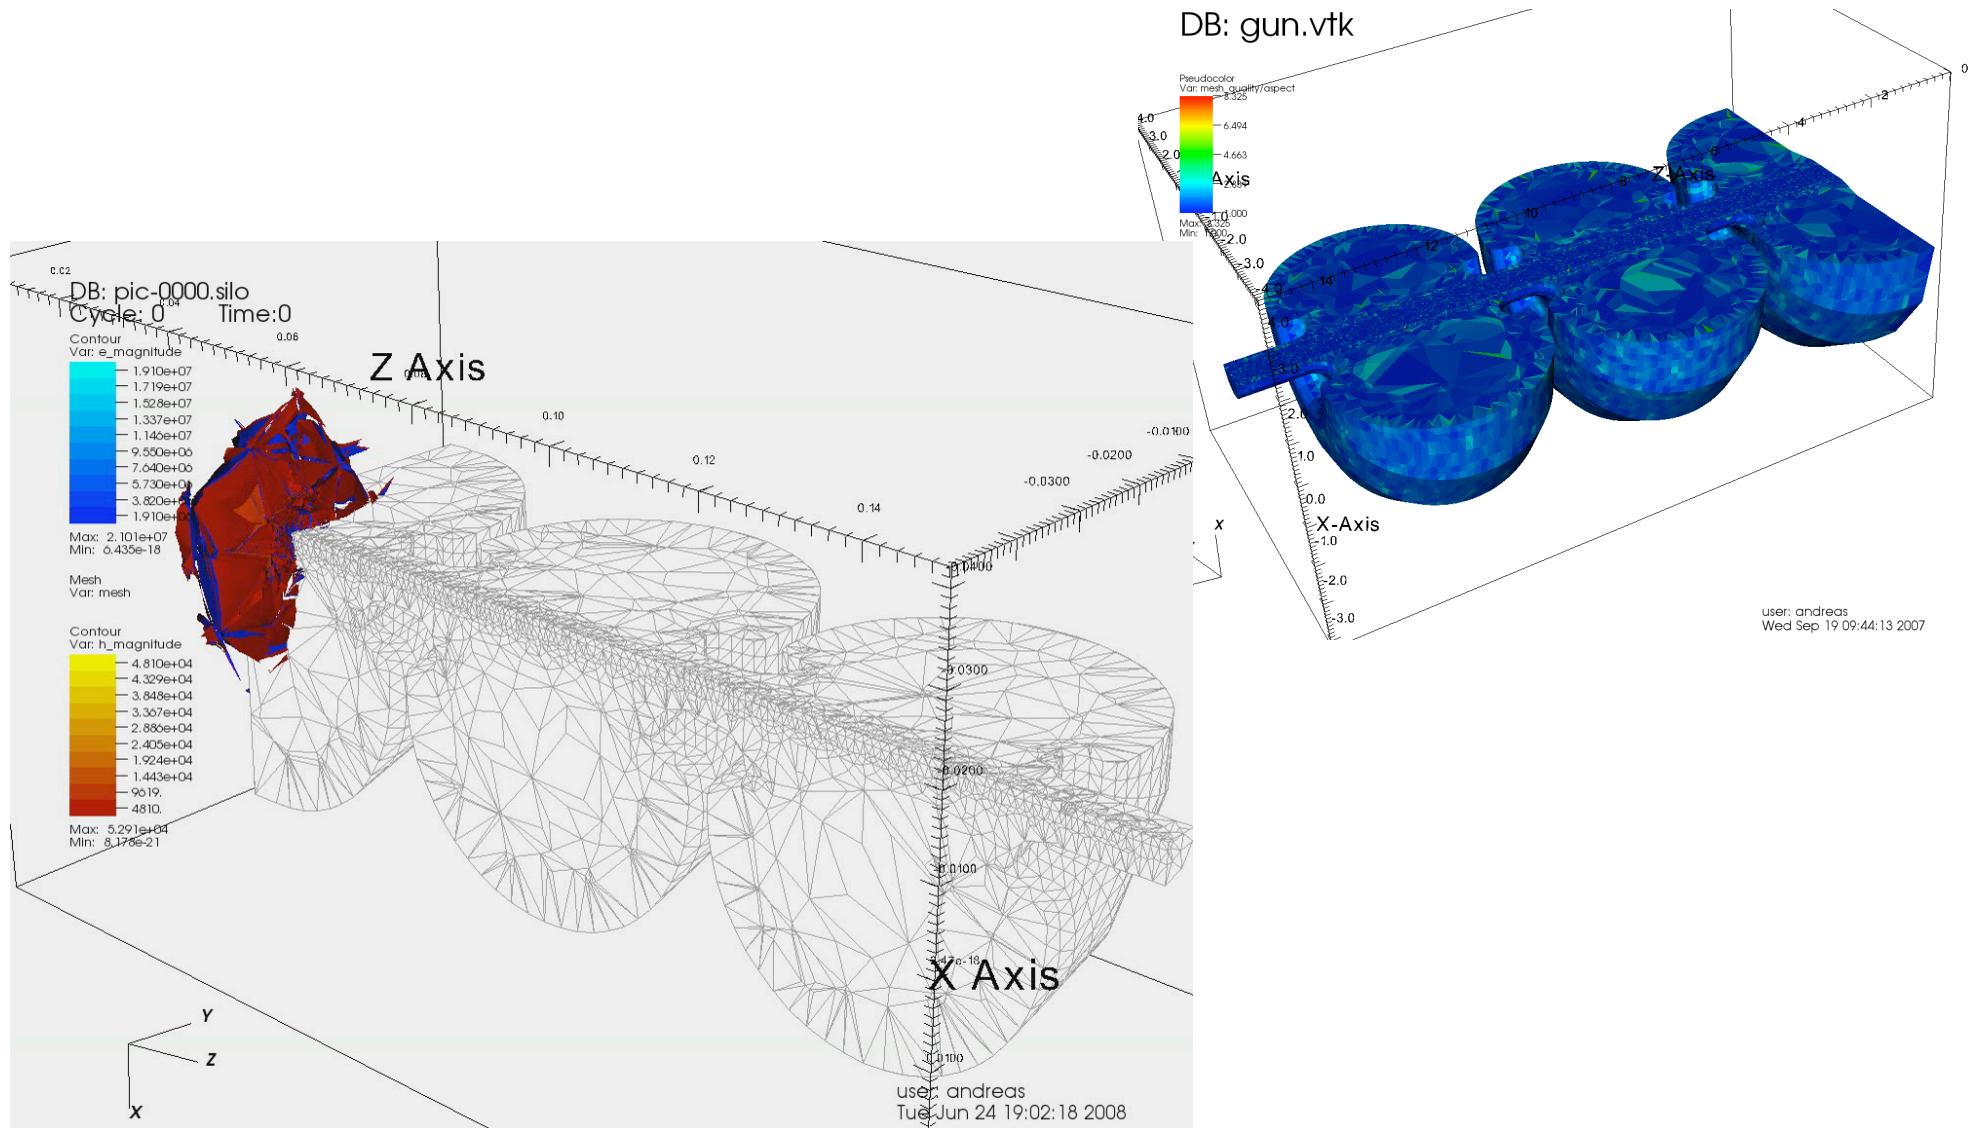
## Fields-to-particles

$$E(x_n), H(x_n)$$

# Some 3D results



# Particle gun



# Summary of Part I

We have generalized everything to 3D

- ✓ Linear/nonlinear problems
- ✓ First order/higher order operators
- ✓ Complex geometries

There is only one significant obstacle to solving large problems

**SPEED !**

# Lecture 8

---

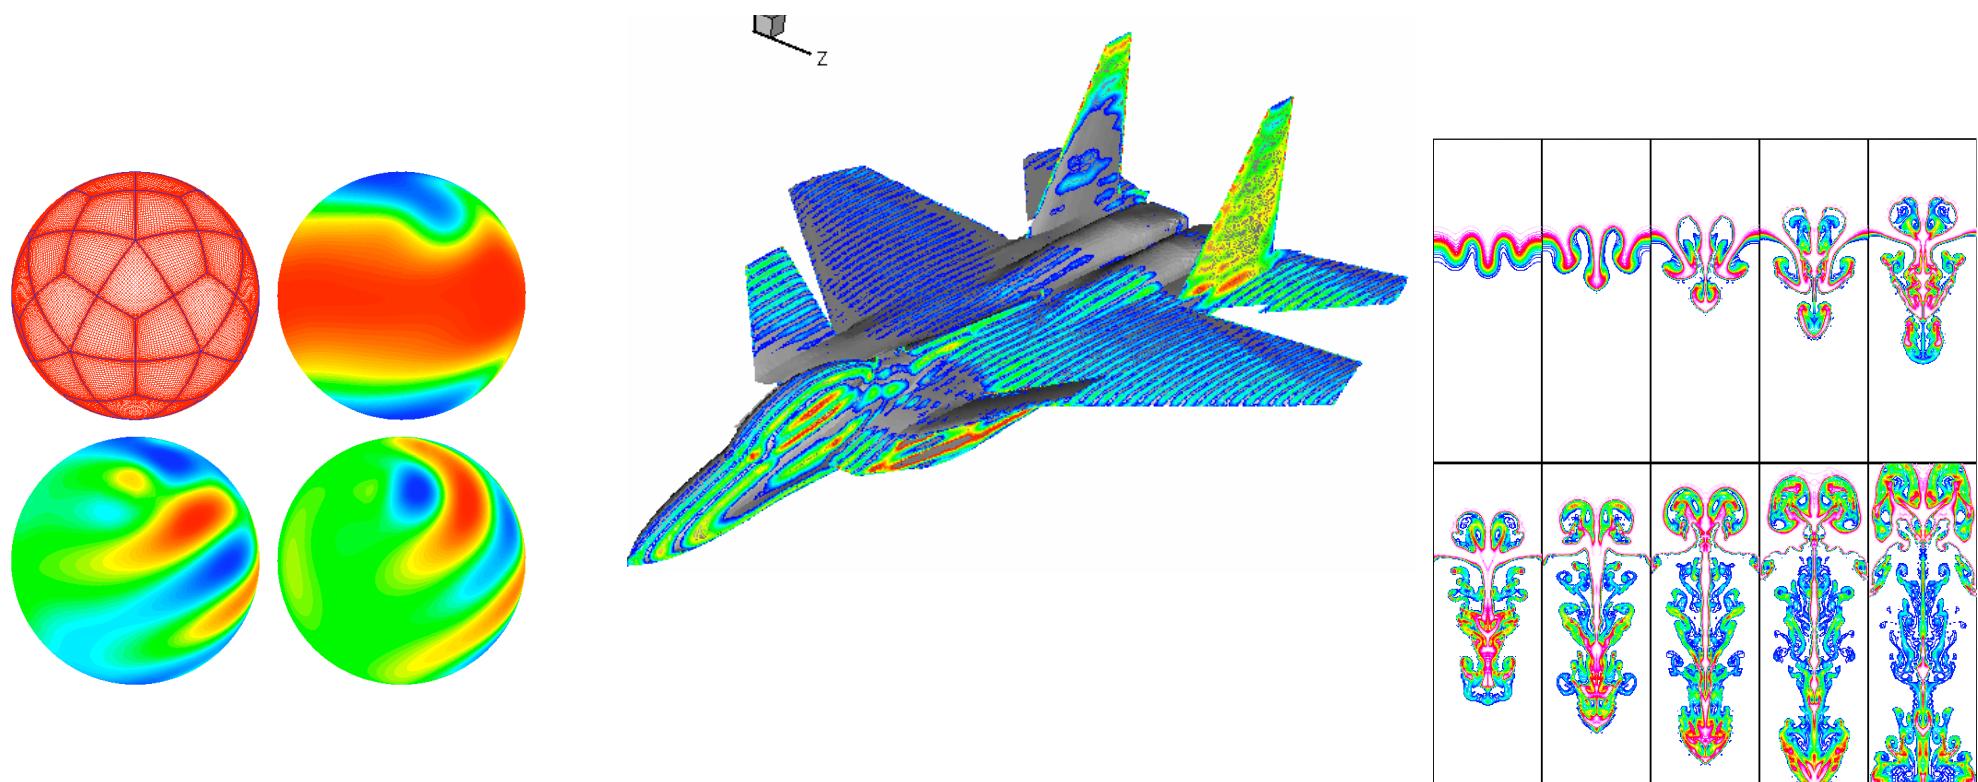
- ✓ Let's briefly recall what we know
- ✓ Part I: 3D problems and extensions
  - ✓ Formulations and examples
  - ✓ Adaptivity and curvilinear elements
- ✓ Part II: The need for speed
  - ✓ Parallel computing
  - ✓ GPU computing
  - ✓ Software beyond Matlab

# The need for speed !

---

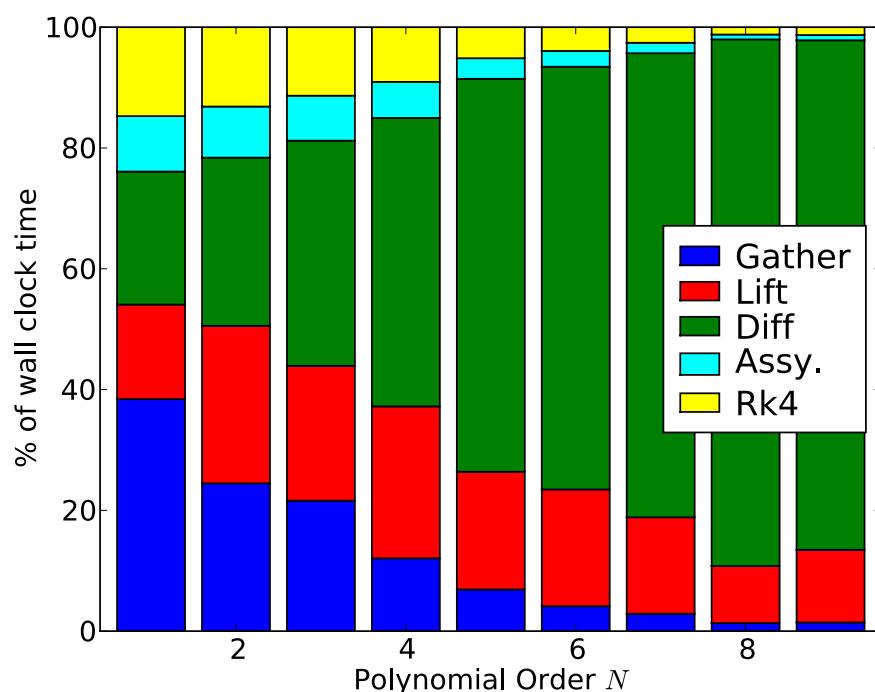
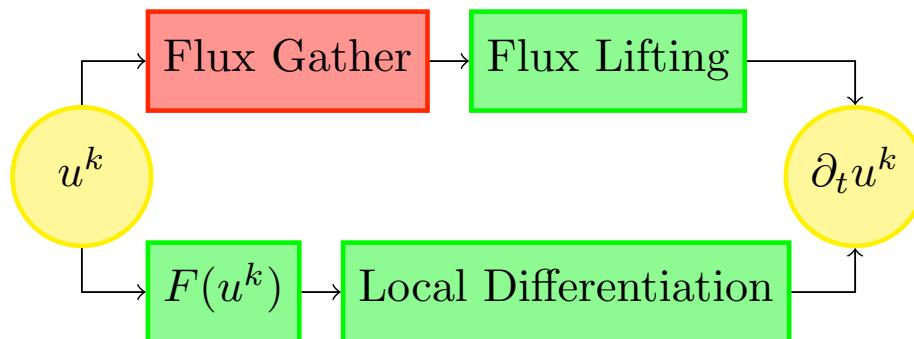
So far, we have focused on ‘simple’ serial computing using Matlab based model.

However, this will not suffice for many applications



# The need for speed

Let us first understand where we spend the time



Test case is  
3D Maxwell's

The majority of  
work is local

# The need for speed

---

The locality suggest that parallel computing will be beneficial

- ✓ Using OpenMP, the local work can be distributed over elements through loops.
- ✓ Using MPI the locality ensures a surface communication model.
- ✓ Mixed OpenMP/MPI models also possible
- ✓ A similar line of arguments can be used for iterative solvers.

# Parallel performance

---

# Processors	64	128	256	512
Scaled RK time	1.00	0.48	0.24	0.14
Ideal time	1.00	0.50	0.25	0.13

High performance is achieved through -

- ✓ Local nature of scheme
- ✓ Pure matrix-matrix operations
- ✓ Local bandwidth minimization
- ✓ Very efficient on-chip performance (~75%)

Challenges -

- ✓ Efficient parallel preconditioning

# Parallel computing

---

DG-FEM maps very well to classic multi-processor computing clusters and result in excellent speed-up.

...but such machines are expensive to buy and run.

**Ex:** To get on the Top500 list, requires about \$3m to purchase a cluster with 16-18Tflop/s performance.

What we need is supercomputing on the desktop

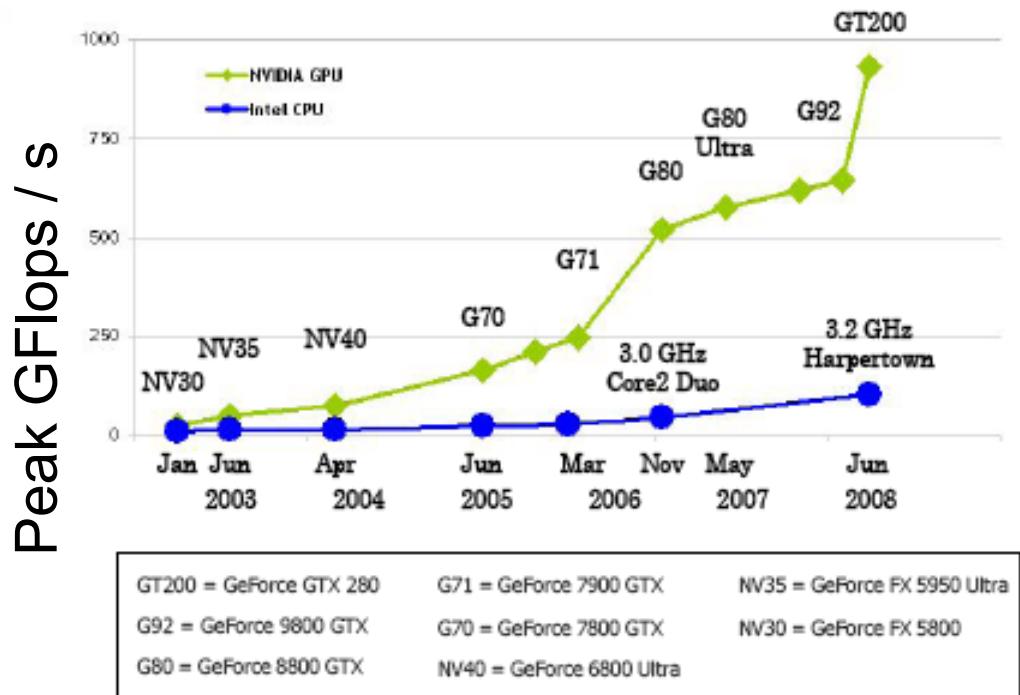
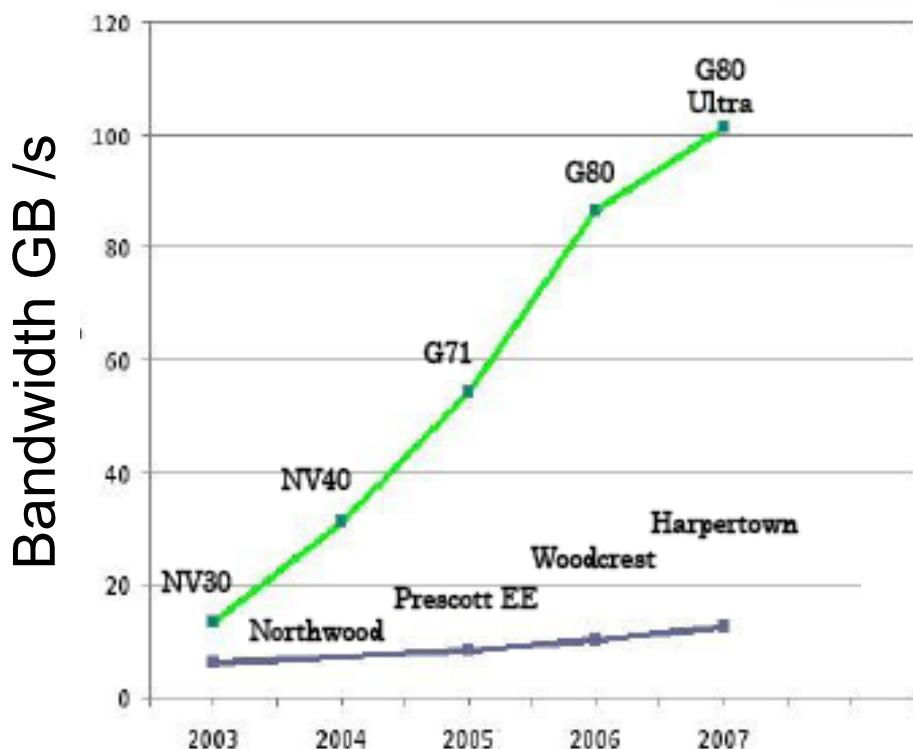
**For FREE !**

... or at least at a fraction of the price

# CPUs vs GPUs



Notice the following



The memory bandwidth and the peak performance on Graphics cards (GPU's) is developing MUCH faster than on CPU's

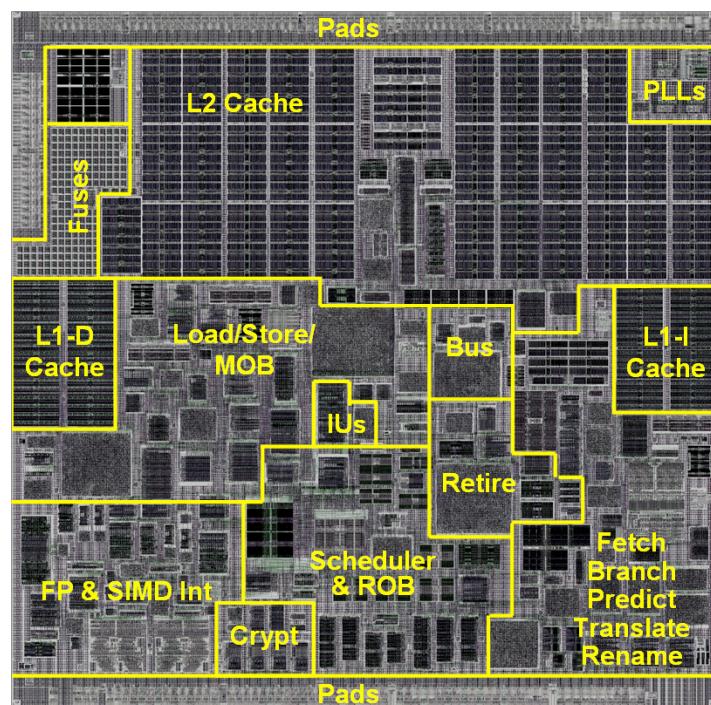
At the same time, the mass-marketing for gaming drives the prices down -- [we have to find a way to exploit this !](#)

# But why is this ?



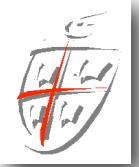
Target for CPU:

- ✓ Single thread very fast
- ✓ Large caches to hide latency
- ✓ Predict, speculate etc



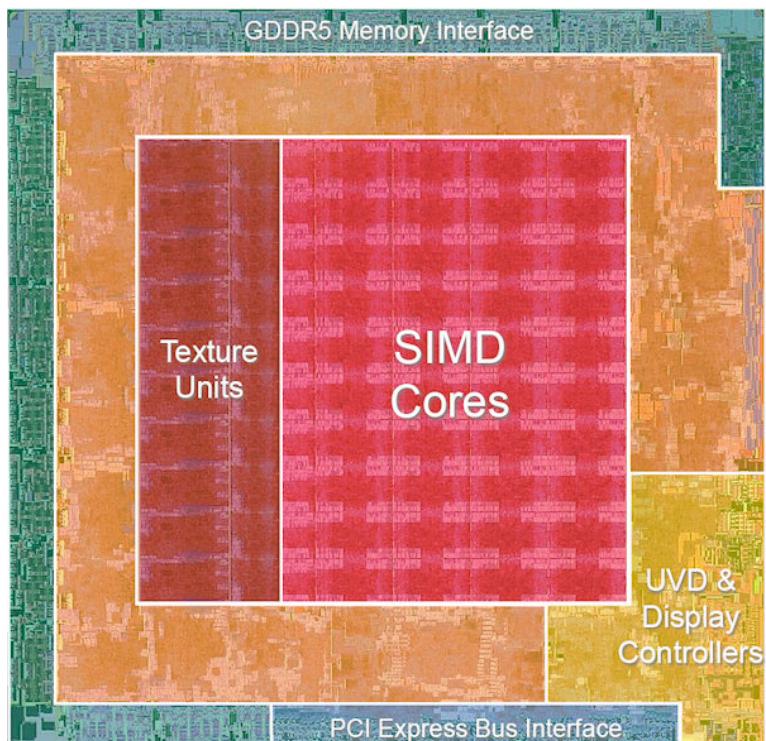
Lots of very complex logic  
to predict behavior

# But why is this ?



For streaming/graphics cards it is different

- ✓ Throughput is what matters
- ✓ Hide latency through parallelism
- ✓ Push hierarchy onto programmer

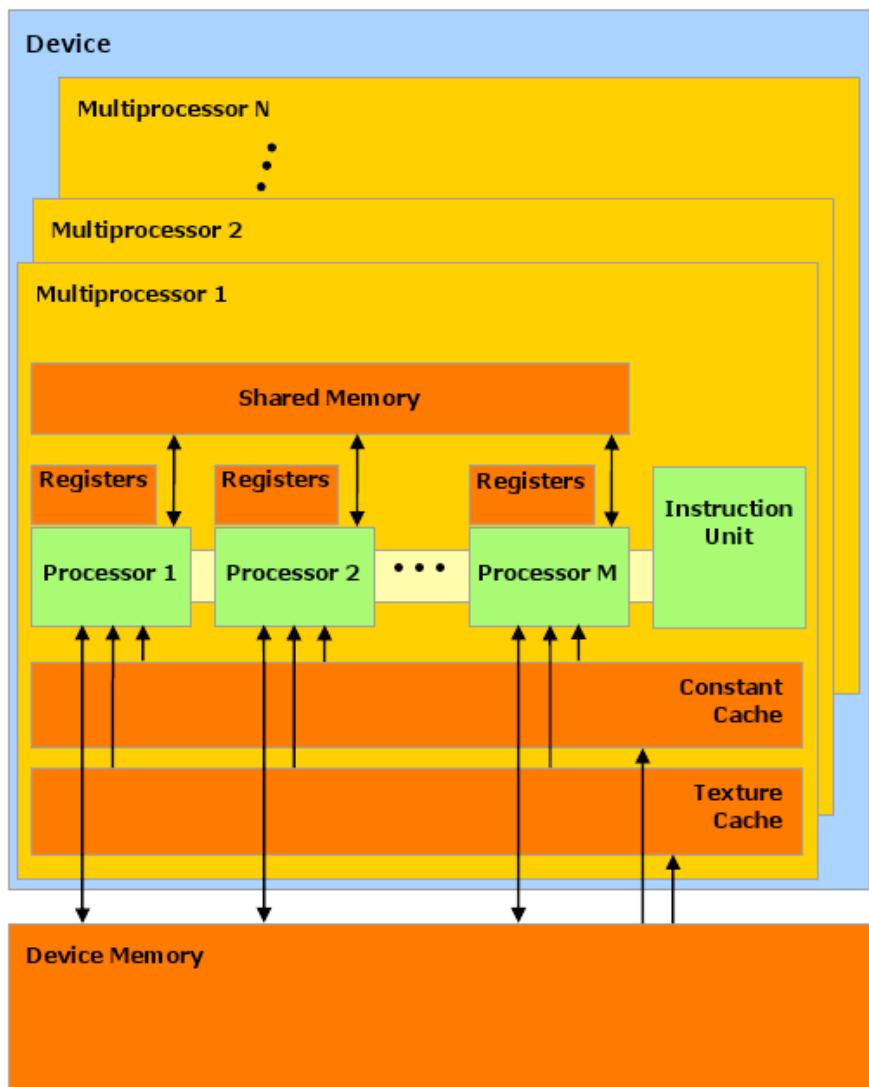


Much simpler logic with  
a focus on performance

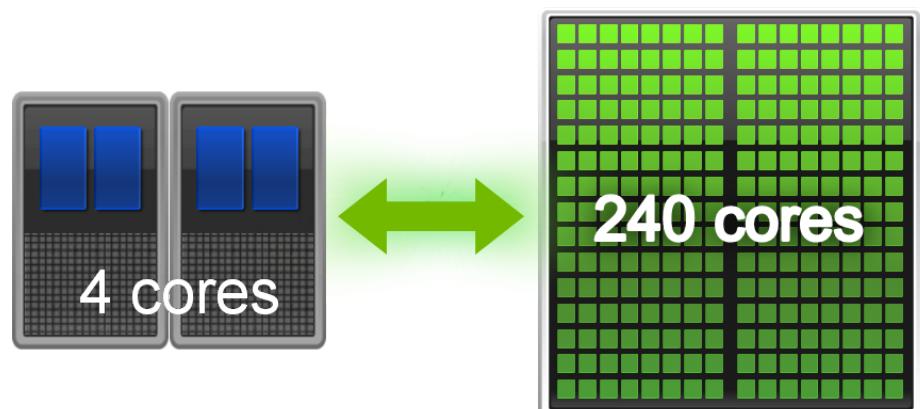
# GPUs 101



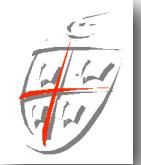
## GPU layout



- ✓ 1 GPU = 30 MPs
- ✓ 1 MP has 1 IU, 8 SP, 1 DP
- ✓ 1 MP has 16KiB shared and 32 KiB Register memory
- ✓ 240 (512) threads
- ✓ Dedicated RAM at 140GB/s
- ✓ Limited caches



# GPUs 101



## Gains

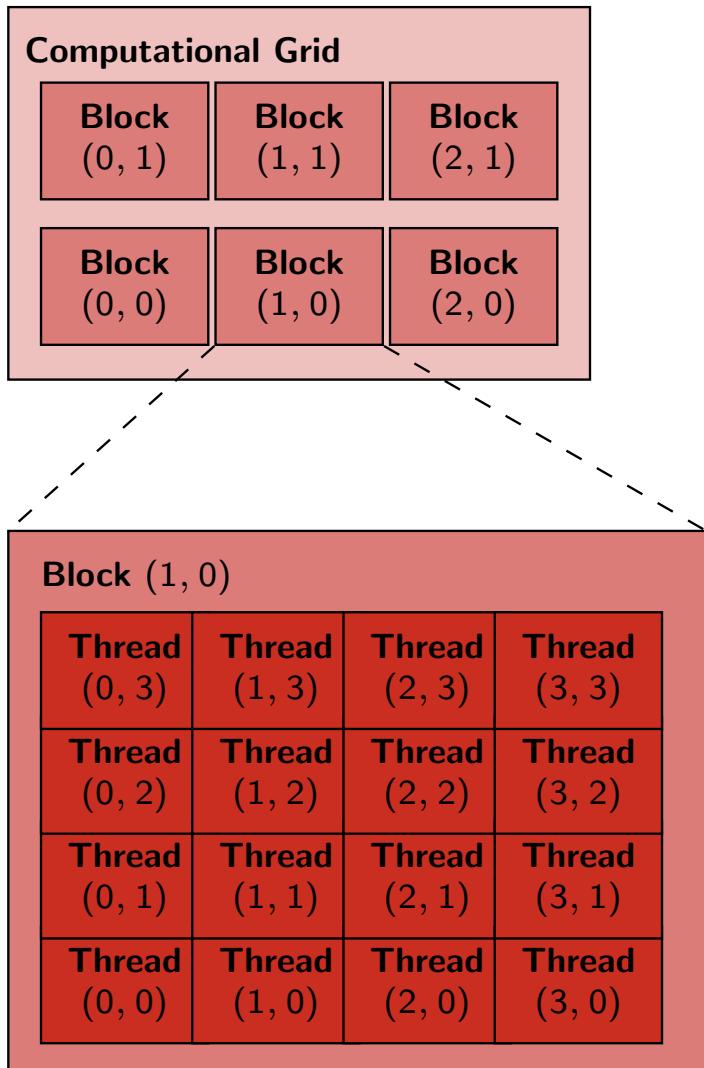
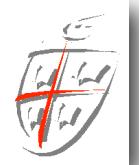
- + Memory Bandwidth  
(140 GB/s vs. 12 GB/s)
- + Compute Bandwidth  
(Peak: 1 TF/s vs. 50 GF/s,  
Real: 200 GF/s vs. 10 GF/s)

## Losses

- Recursion
- Function pointers
- Exceptions
- IEEE 754 FP compliance
- Cheap branches (i.e. ifs)

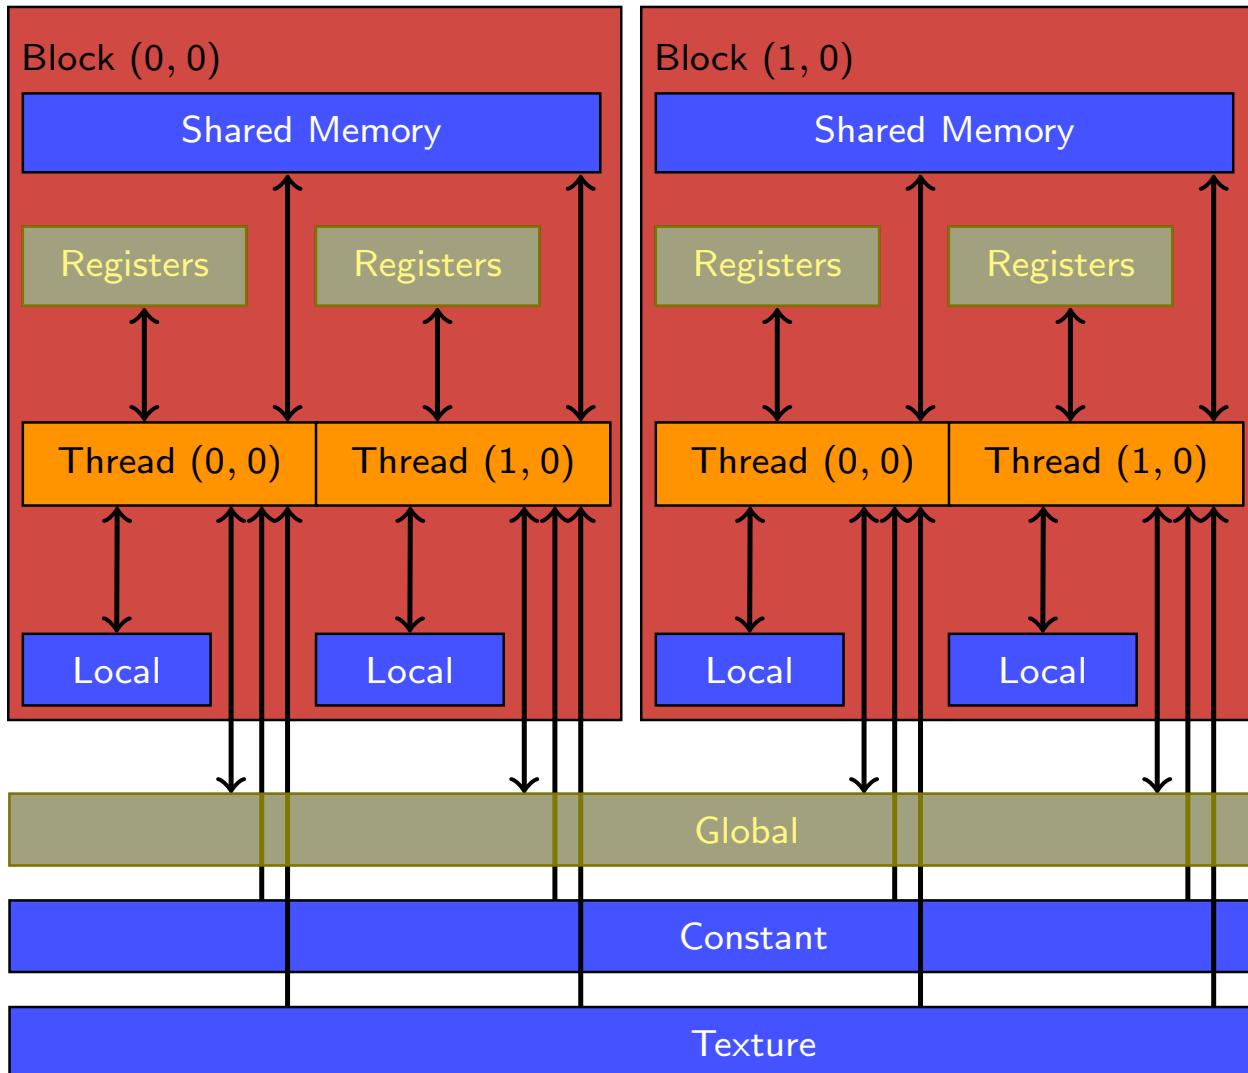
Already here it is clear that programming models/codes may have to undergo substantial changes -- and that not all will work well

# GPUs 101



- ✓ Genuine multi-tiered parallelism
  - ✓ Grids
  - ✓ blocks
  - ✓ threads
- ✓ Only threads within a block can talk
  - ✓ Blocks must be executed in order
- ✓ Grids/blocks/threads replace loops
- ✓ Until recently, only single precision
- ✓ Code-able with CUDA (C-extension)

# GPUs 101



Memory model:

- ✓ Registers
- ✓ Local shared
- ✓ Global

# GPUs 101

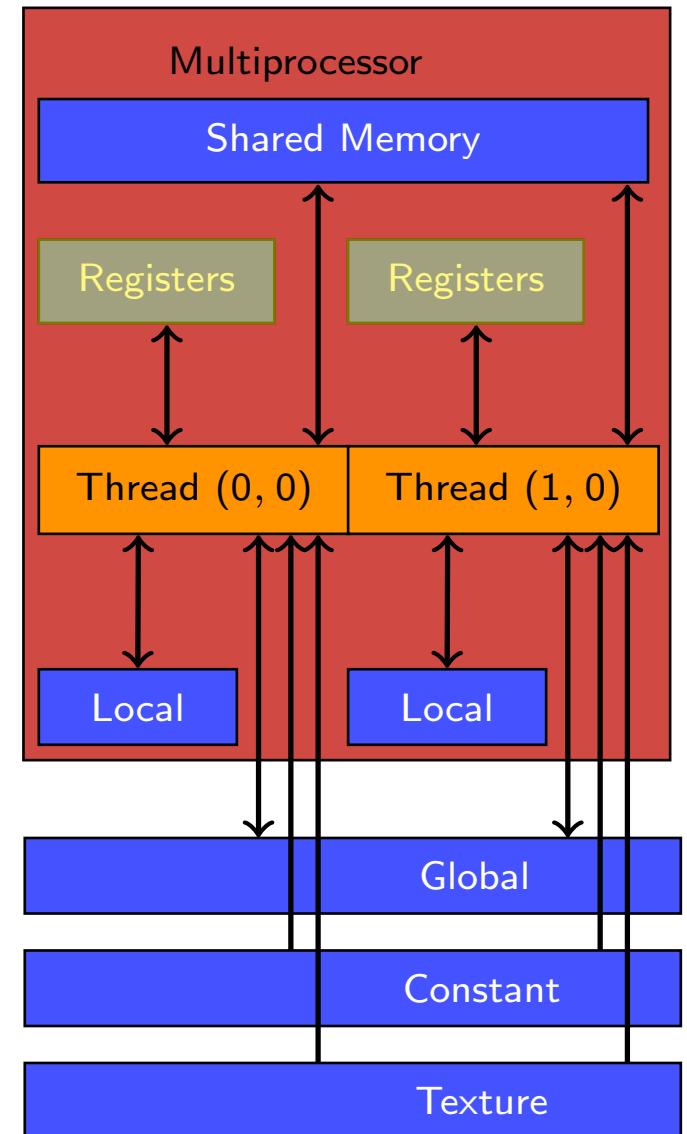


✓ Lots of multi-processors (about 30)

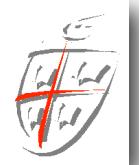
... communicate through global mem

✓ Registers, shared memory, and threads communicate with low latency

... but memory is limited (16-32 KiB)



# GPUs 101



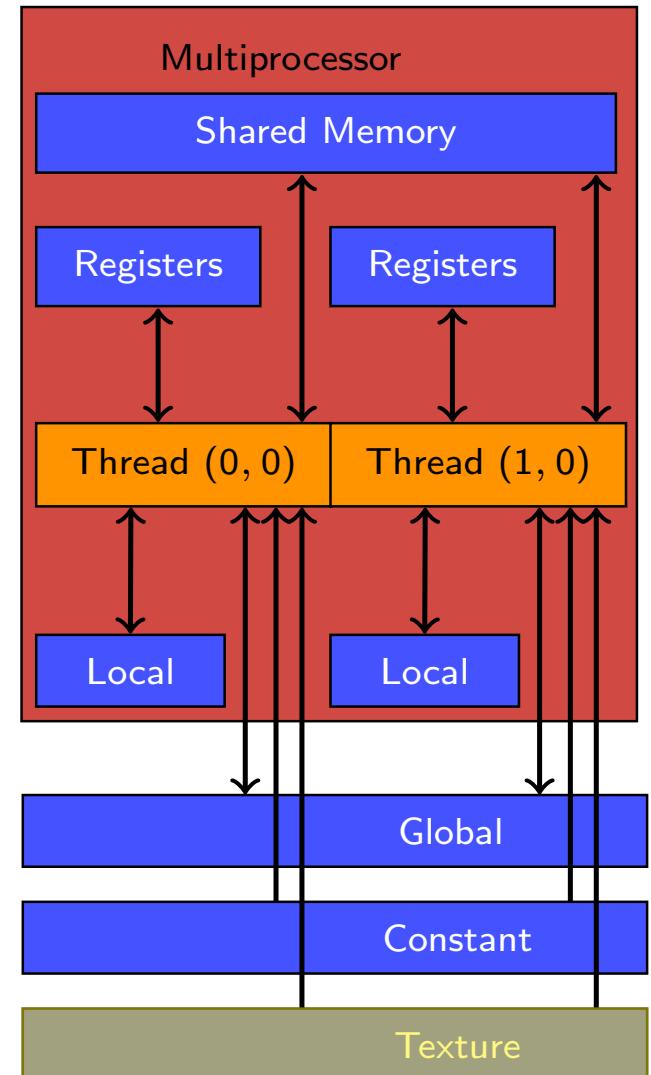
✓ Global memory (4GiB/PU) is plentiful

... but latency is high (512 bit bus)  
... and stride one is preferred

✓ Texture is similar to global memory

... allows more general access patterns  
... but it is read only

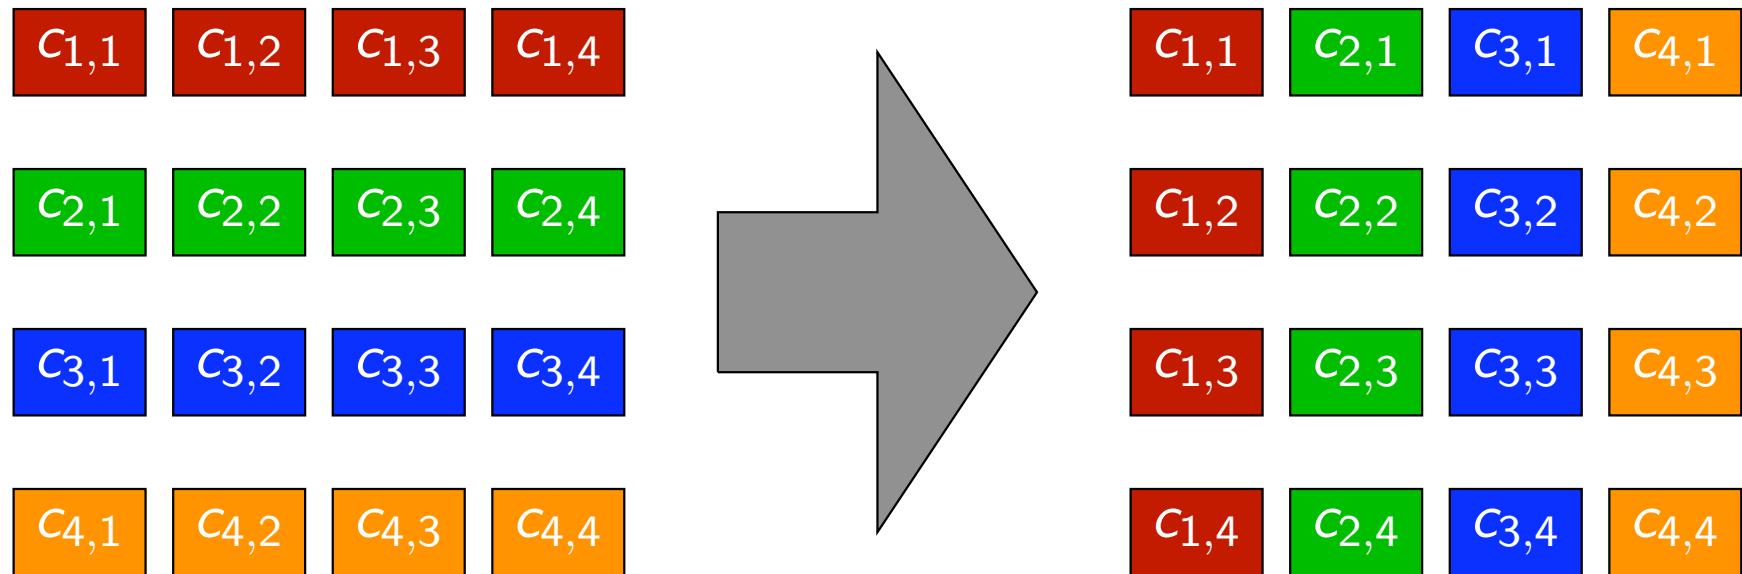
Type	Per	Access	Latency
Registers	thread	R/W	1
Local	thread	R/W	1000
Shared	block	R/W	2
Global	grid	R/W	1000
Constant	grid	R/O	1-1000
Texture	grid	R/O	1000



# Let's consider an example



## Matrix transpose

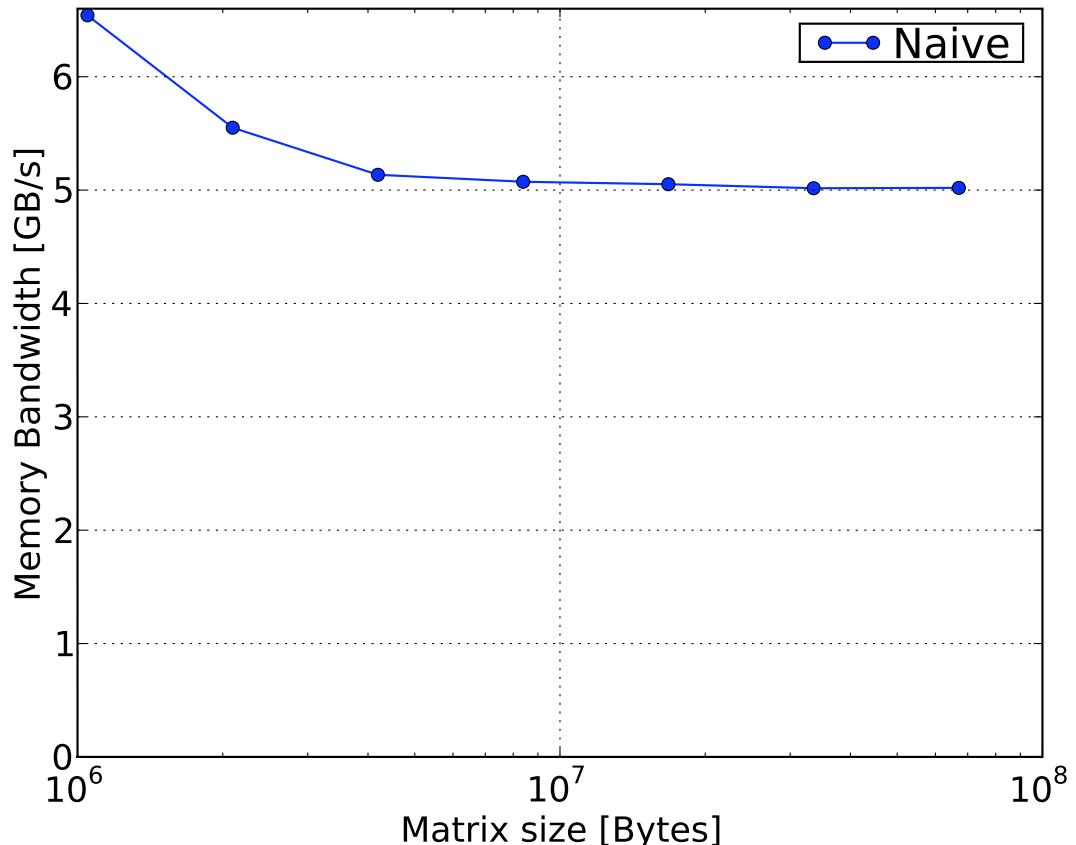


Memory bandwidth will be a limit here

# Let's consider an example

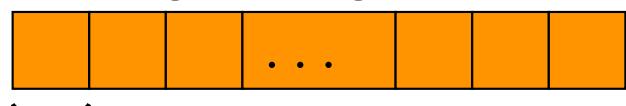


## Using just global memory



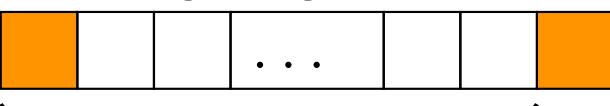
As CPU

Reading from global mem:



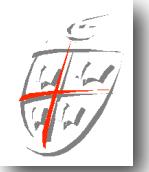
stride: 1 → one mem.trans.

Writing to global mem:

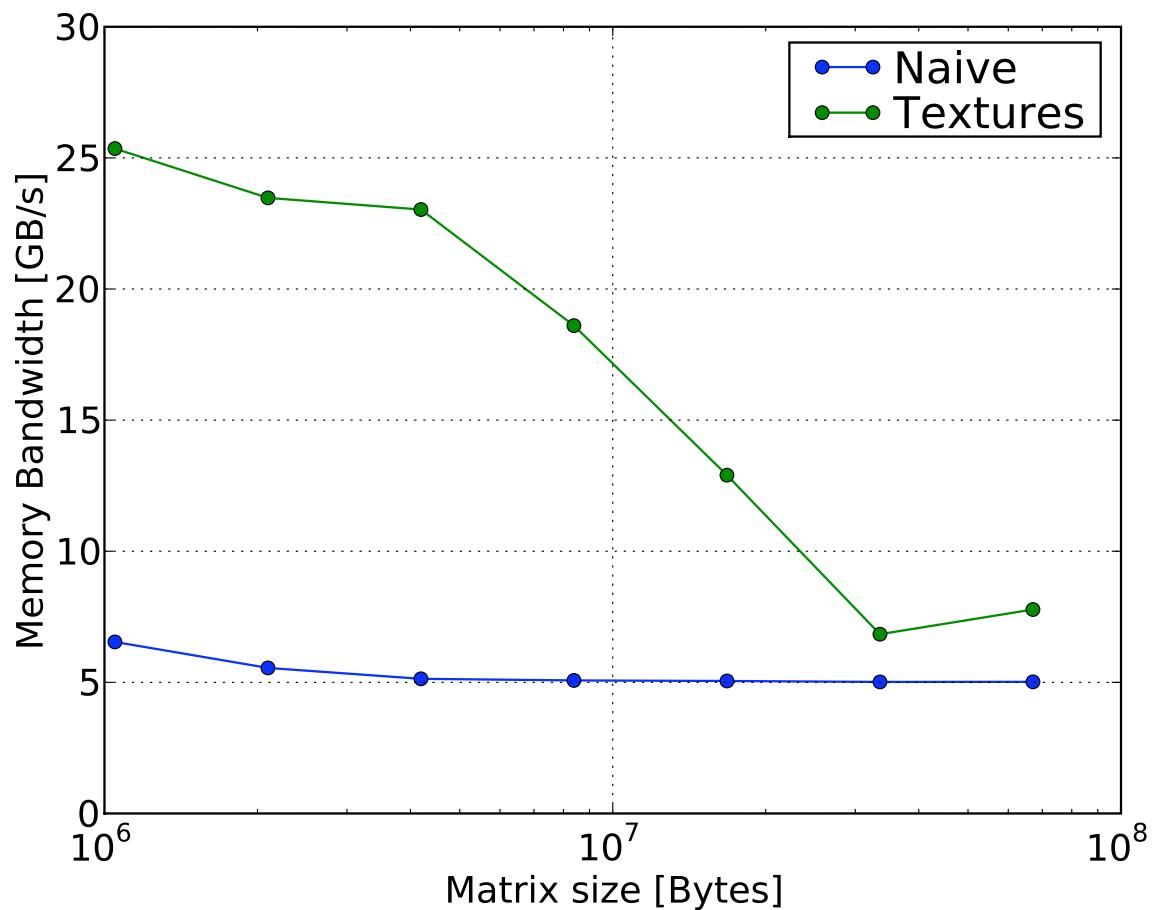


stride: 16 → 16 mem.trans.!

# Let's consider an example



Using just texture(read)+global(write) memory

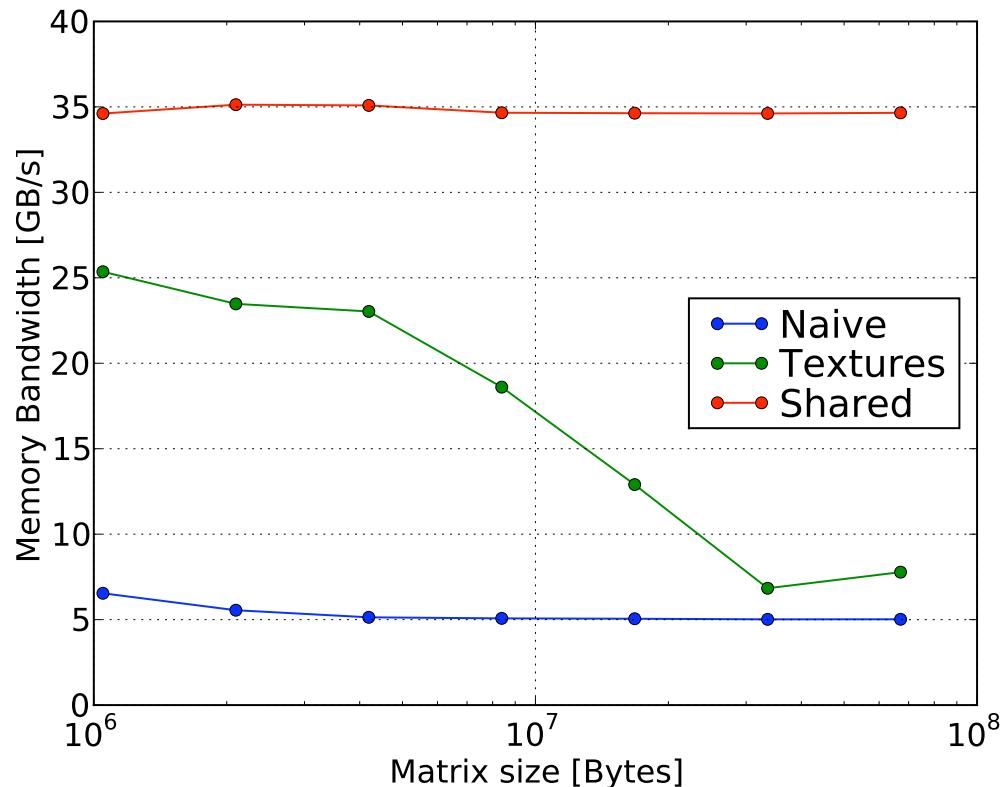


Getting better

# Let's consider an example



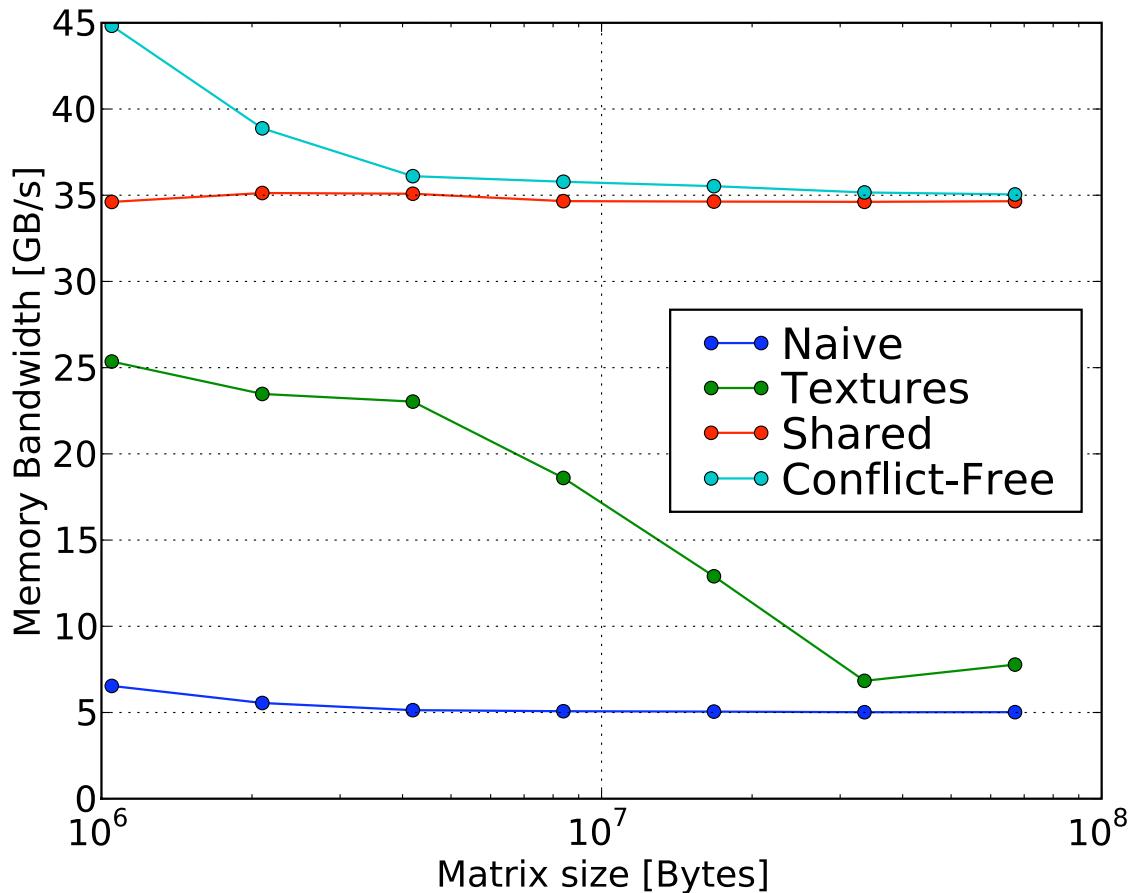
Transpose block-by-block in shared memory -  
this does not care about strides



# Let's consider an example



Additional improvements are possible for small matrices - bank conflicts in shared memory



A factor of 7-8 over CPU

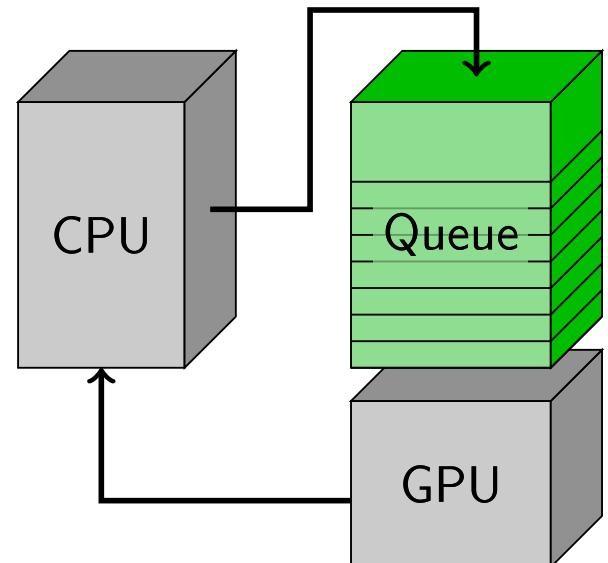
# CPUs vs GPUs



As T.Warburton said in his talk:

*The CPU is mainly the traffic controller  
... although it need not be*

- ✓ The CPU and GPU runs asynchronously
- ✓ CPU submits to GPU queue
- ✓ CPU synchronizes GPUs
- ✓ Explicitly controlled concurrency is possible



# GPUs overview

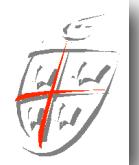
---



- ✓ GPUs exploit multi-layer concurrency
  - ✓ The memory hierarchy is deep
  - ✓ Memory padding is often needed to get optimal performance
  - ✓ Several types of memory must be used for performance
- 

- ✓ First factor of 5 is not too hard to get
- ✓ Next factor of 5 requires quite some work
- ✓ Additional factor of 2-3 requires serious work

# Nodal DG on GPU's

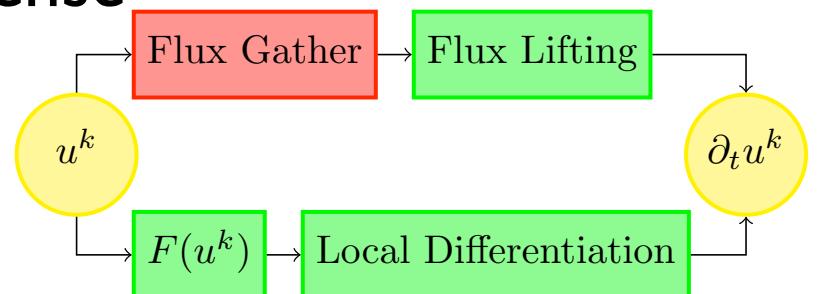


So what does all this mean ?

✓ GPU's has deep memory hierarchies so local is good  
→ The majority of DG operations are local

✓ Compute bandwidth >> memory bandwidth  
→ High-order DG is arithmetically intense

✓ GPU global memory favors dense data  
→ Local DG operators are all dense

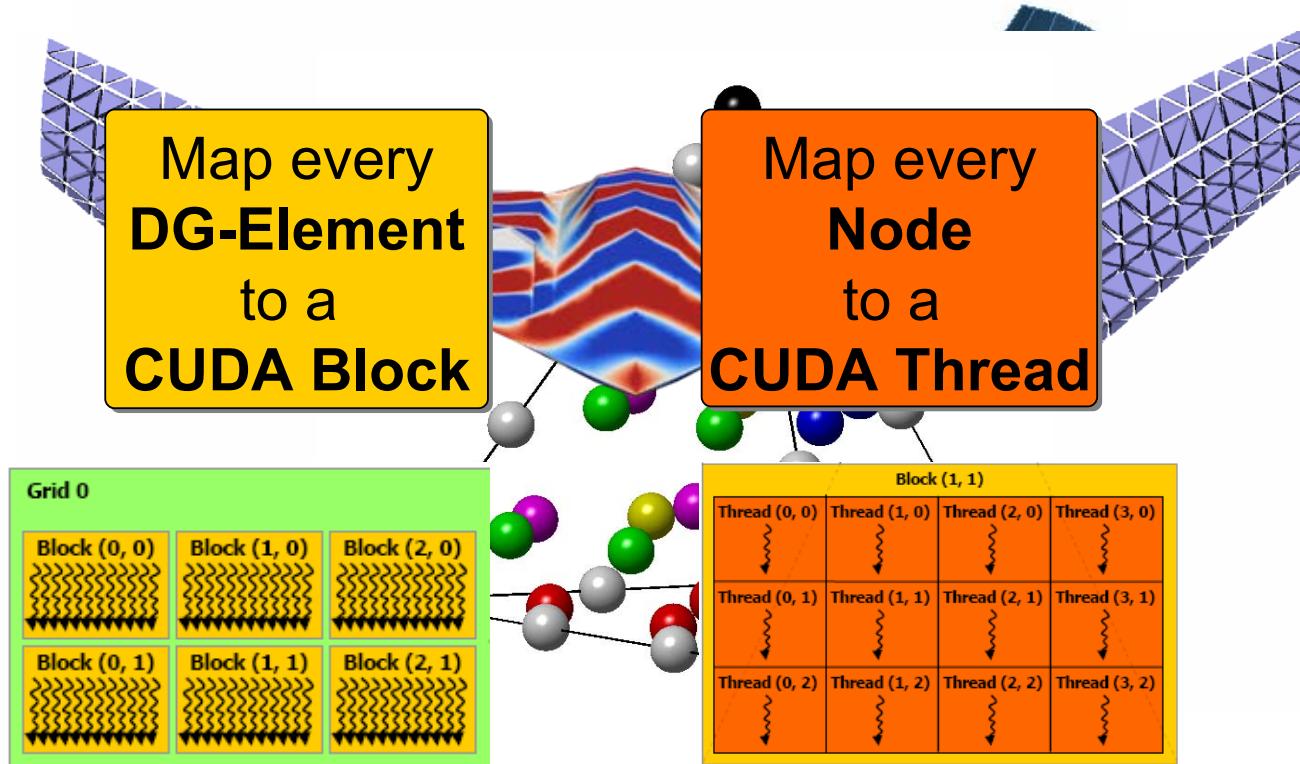


With proper care we should be able to obtain excellent performance for DG-FEM on GPU's

# Nodal DG on GPU's



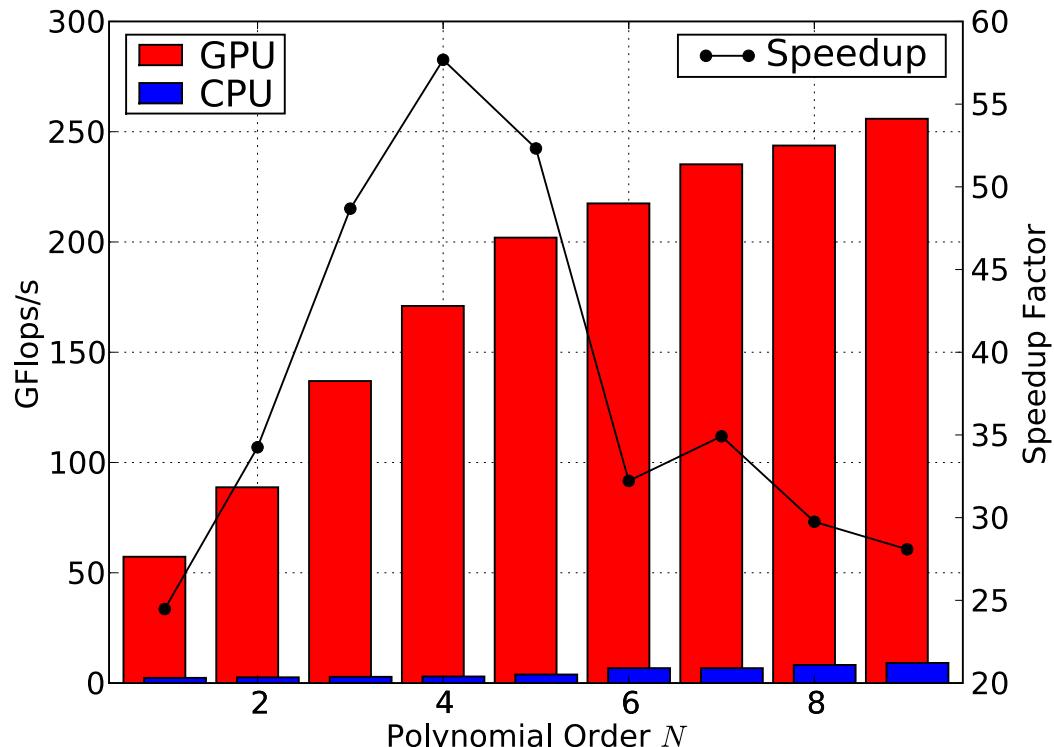
Nodes in threads, elements for blocks



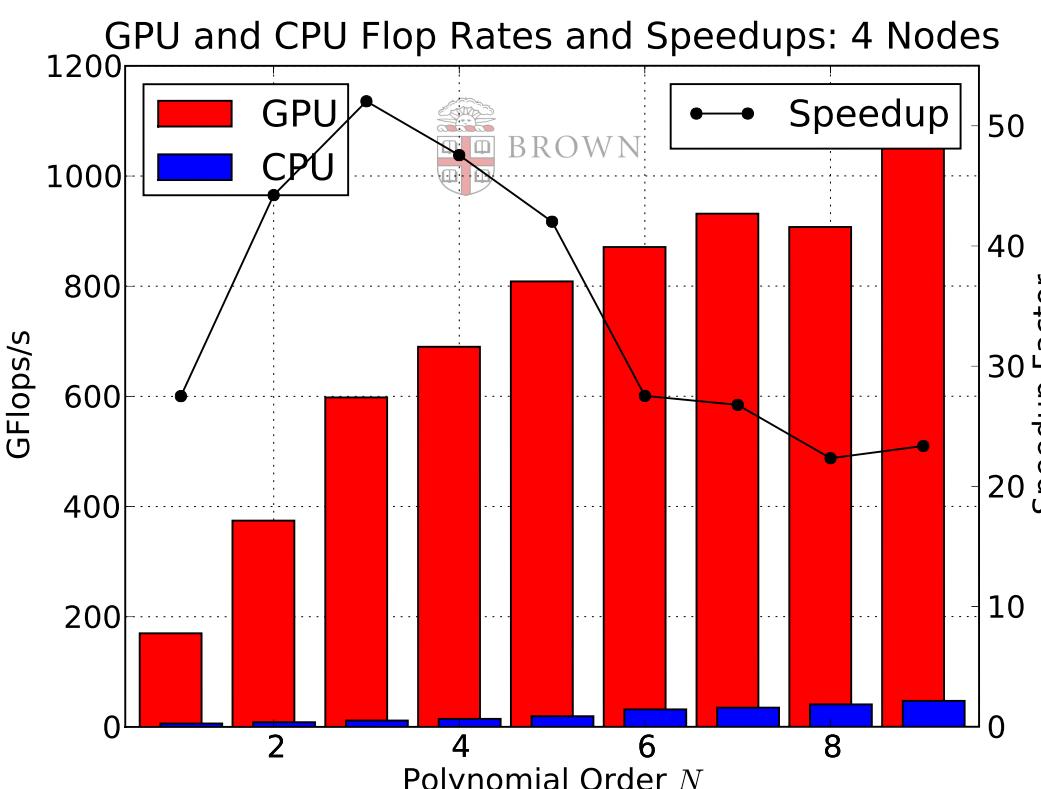
Other choices:

- ✓ D-matrix in shared, data in global (small N)
- ✓ Data in shared, D-matrix is global (large N)

# Nodal DG on GPU's



DG-FEM on one GPU

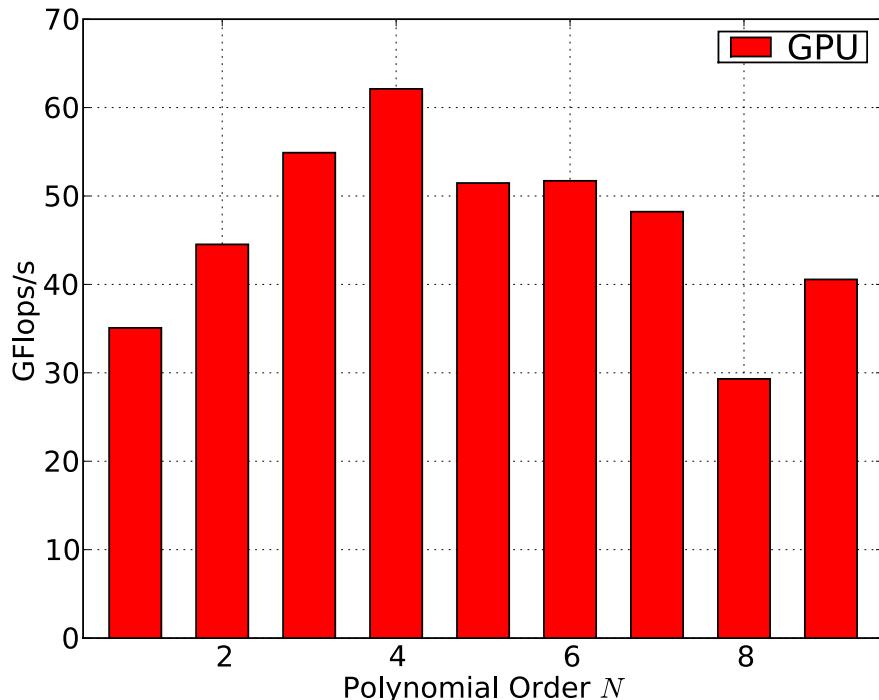


DG-FEM on four GPU  
one card

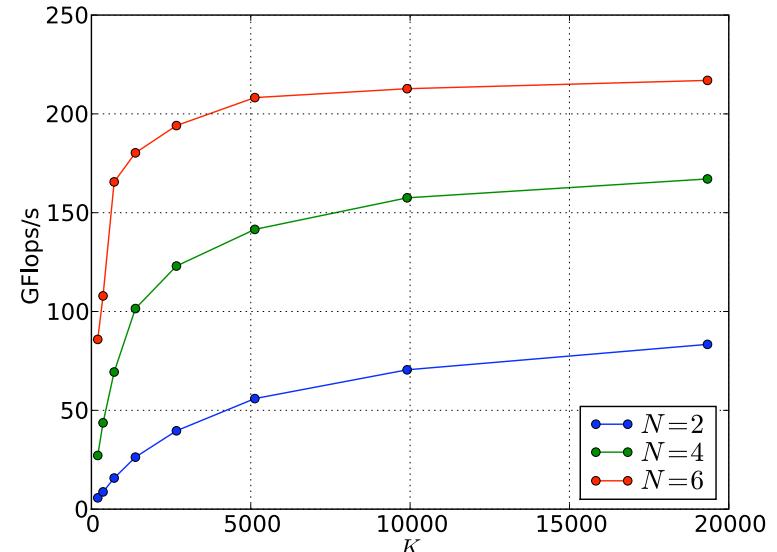
# Nodal DG on GPU's



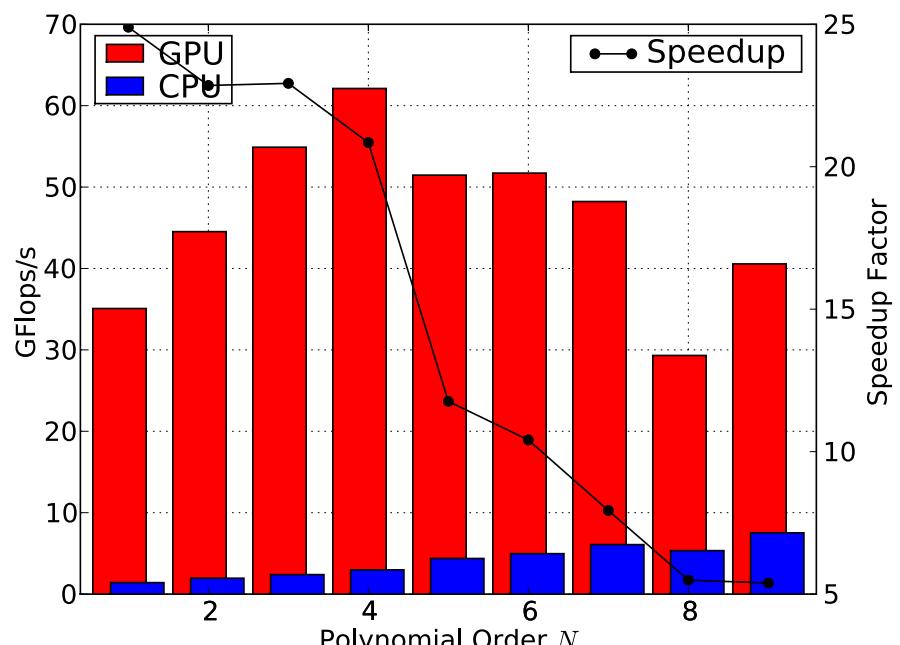
Where you need it most



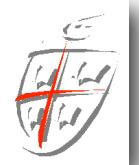
... and for larger and larger grids



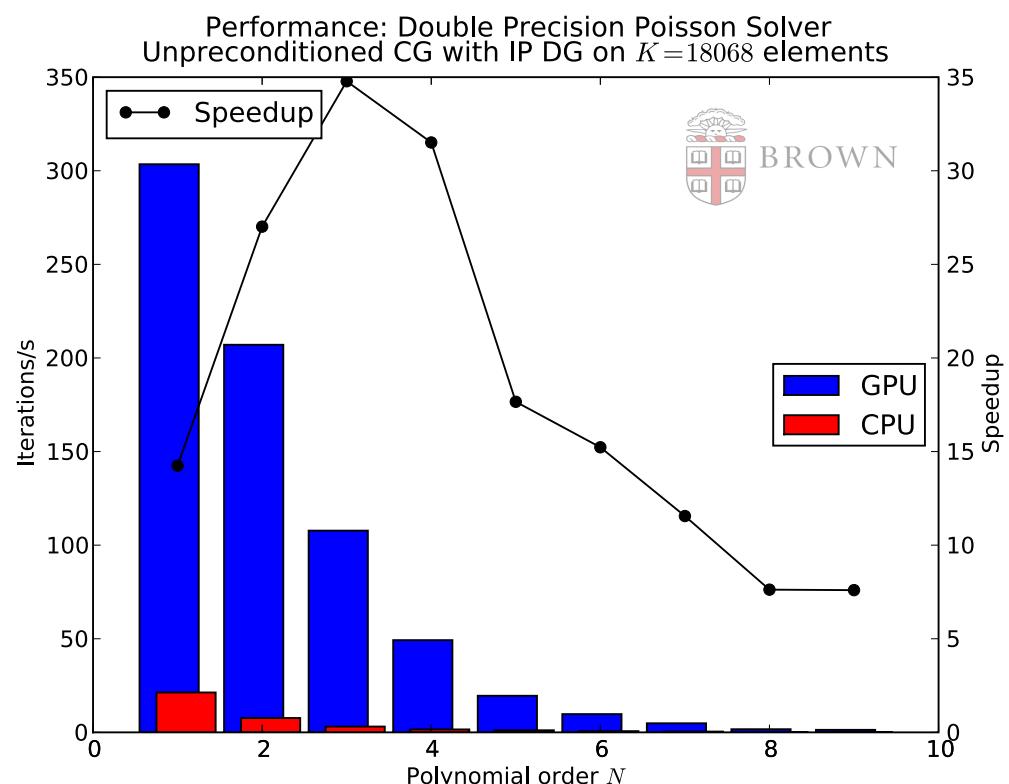
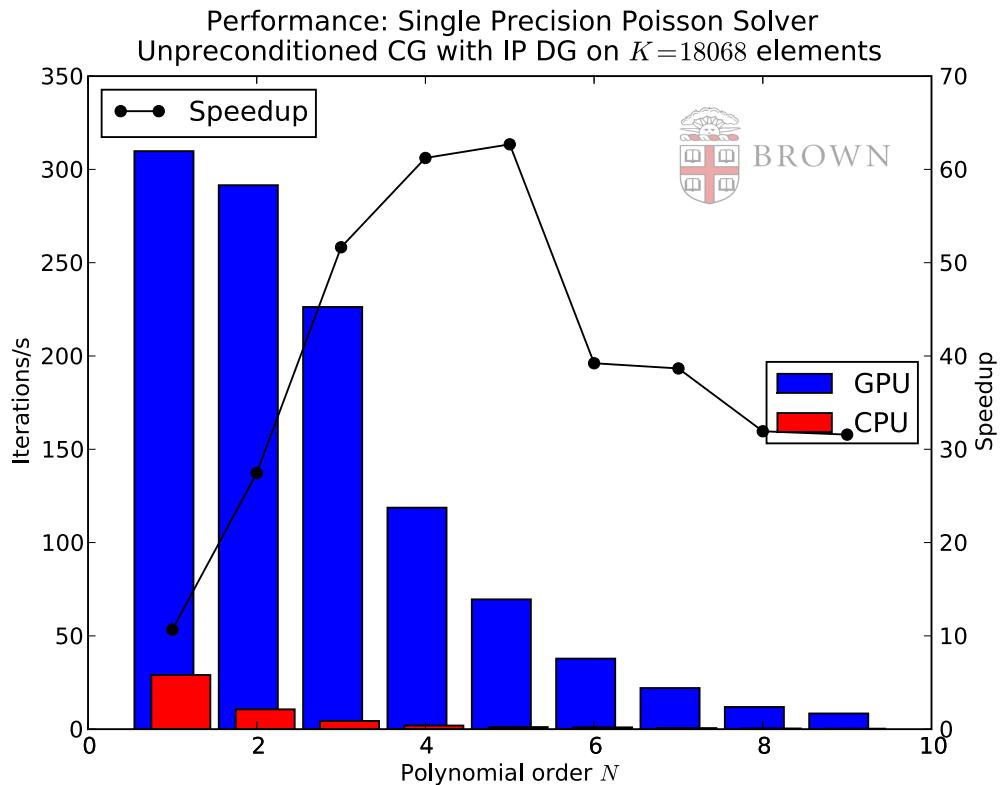
Also in double precision



# Nodal DG on GPU's



## Similar results for DG-FEM Poisson solver with CG



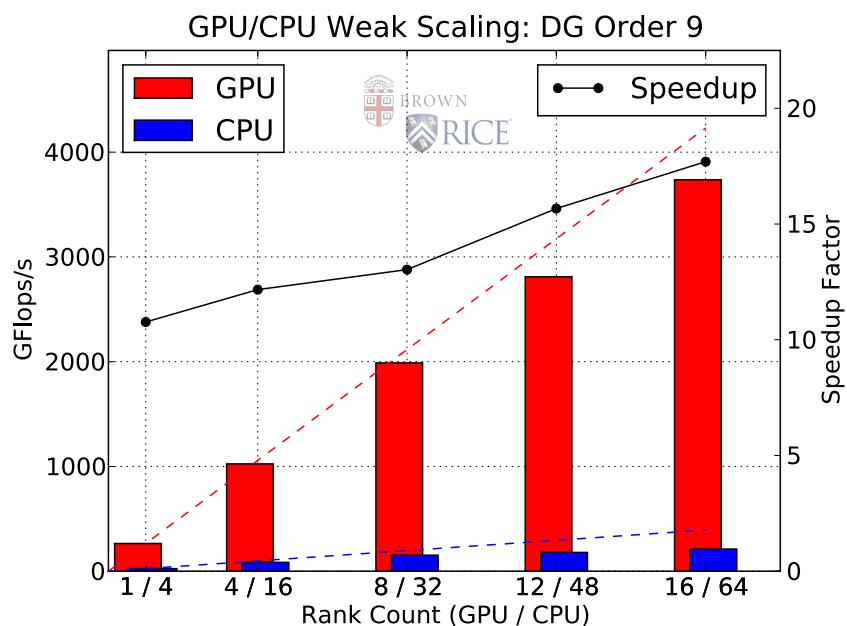
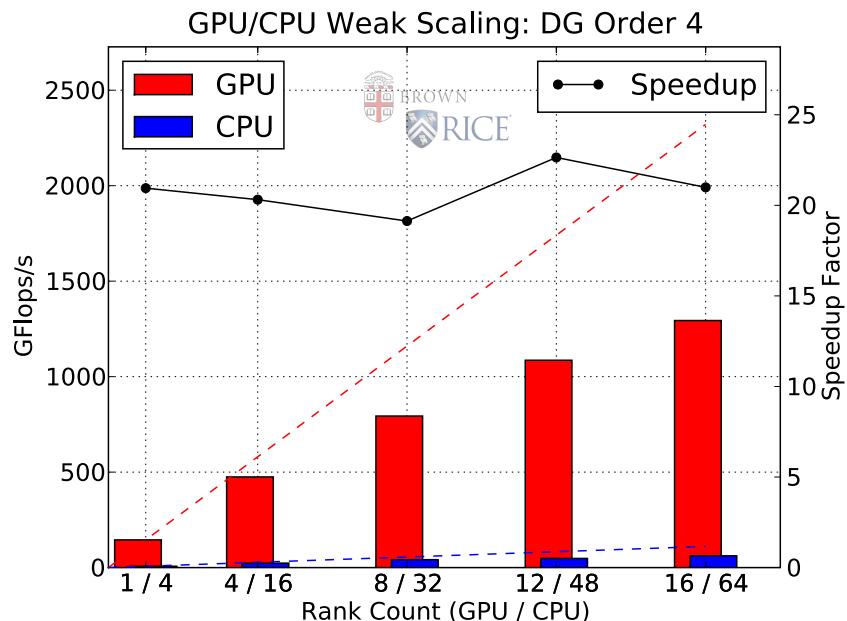
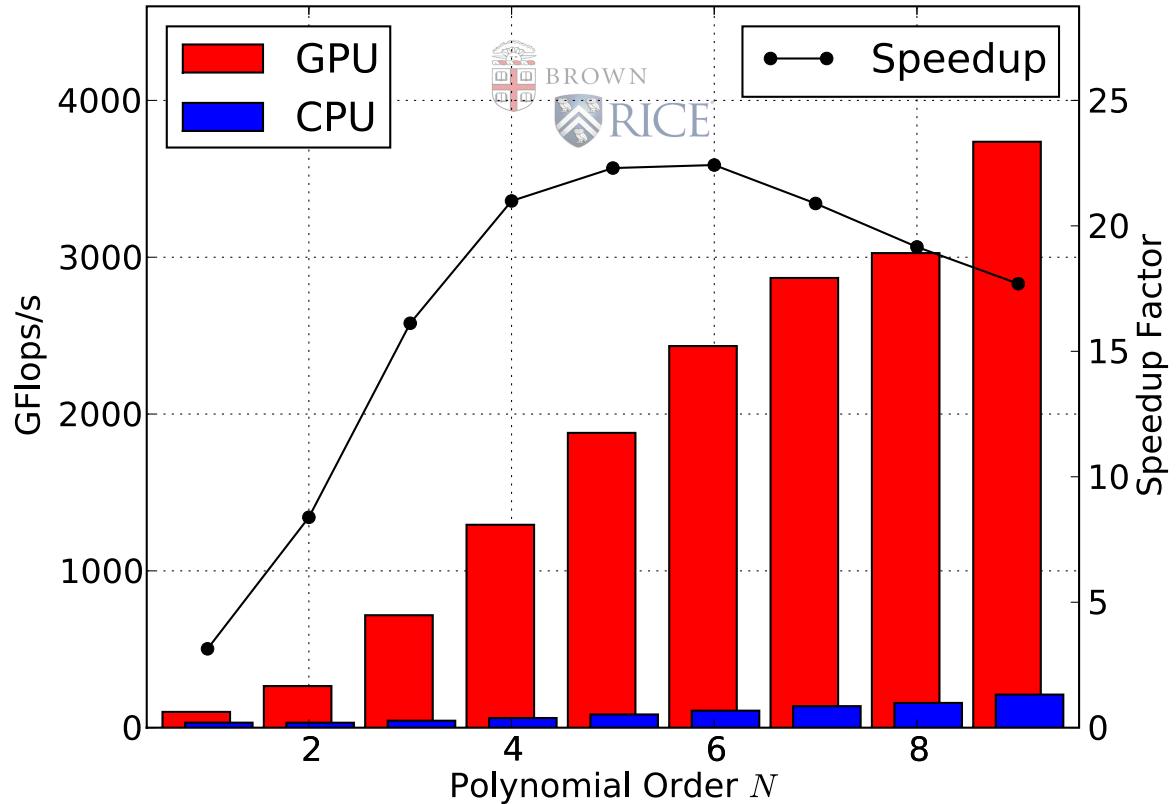
Note: No preconditioning

# Combined GPU/MPI solution



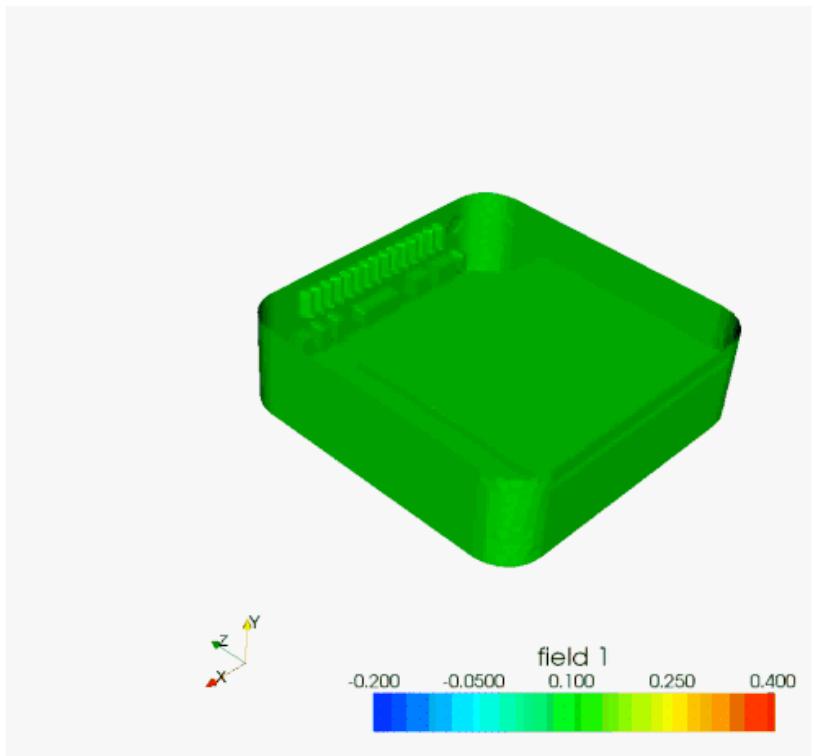
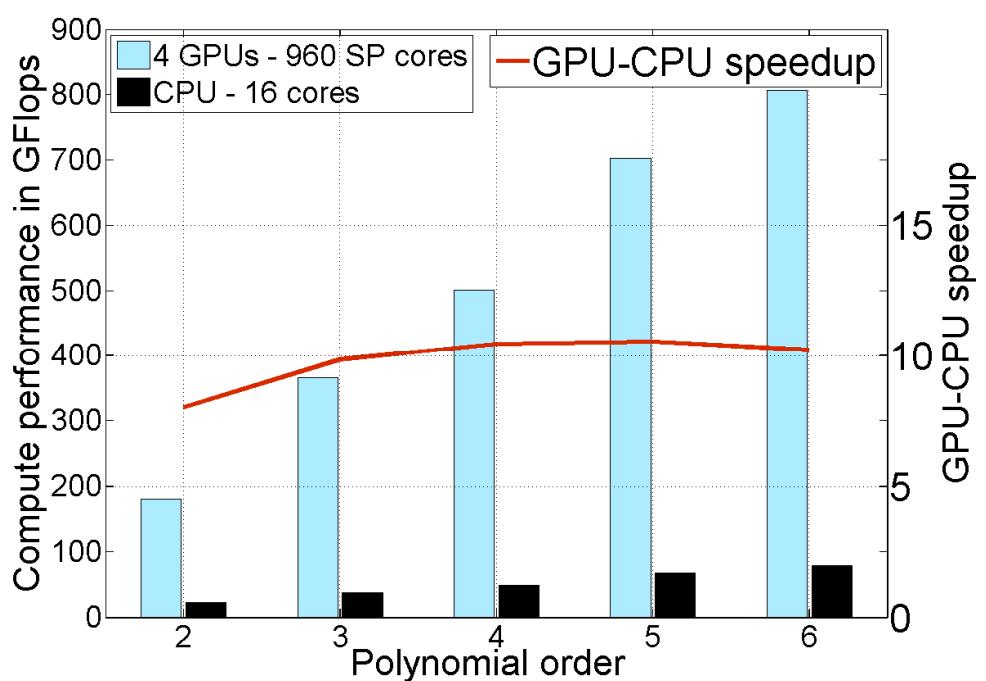
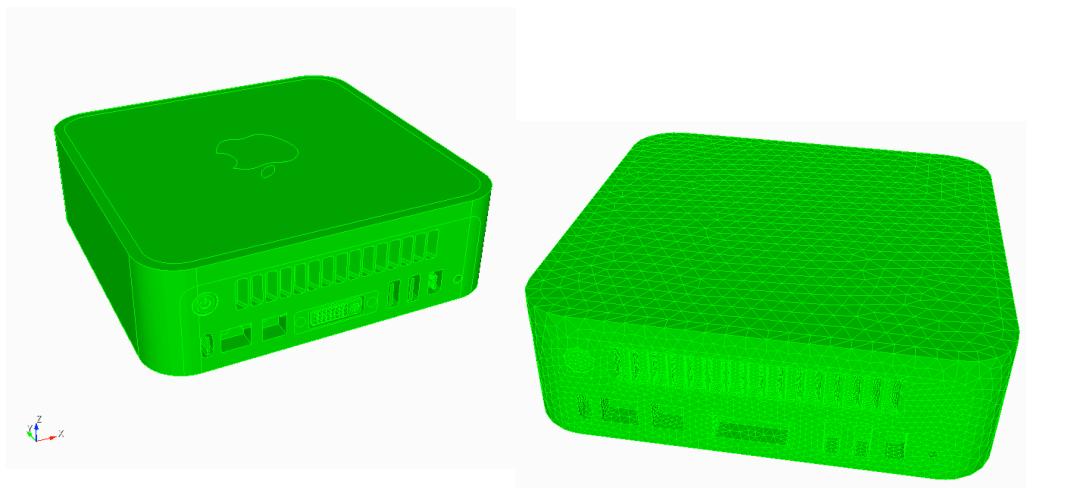
## MPI across network

Flop Rates and Speedups: 16 GPUs vs 64 CPU cores



Good scaling when problem  
is large

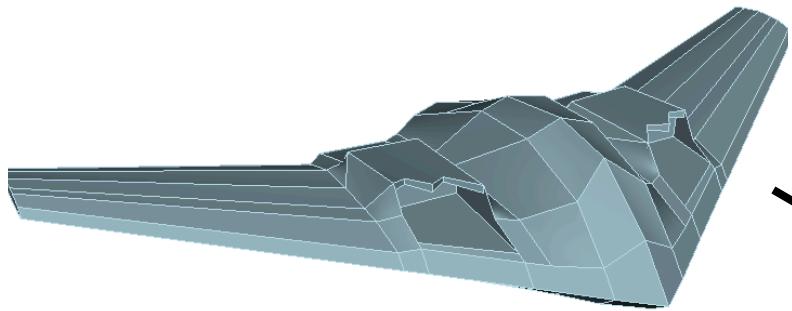
# Example - a Mac Mini



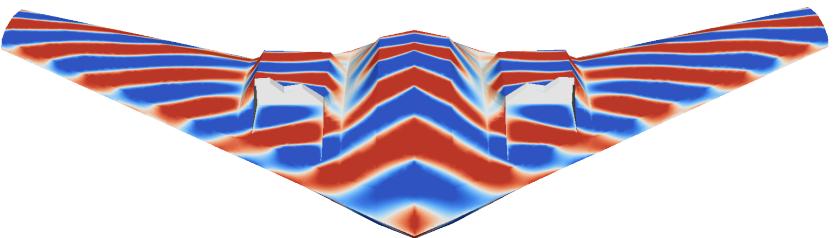
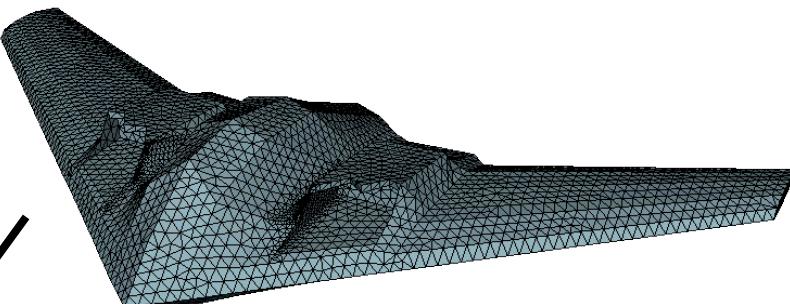
K=201765 elements  
3rd order elements

Computation by N. Godel

# Example: Military aircraft



$K=130413$  elements  
3rd order elements  
 $15.6E6$  DoF



Computation by N. Godel

	CPU global	29 h 6 min 46 s	1.0
	GPU global	39 min 1 s	44.8
	GPU multirate	11 min 50 s	147.6

# Nodal DG on GPU's



Not just for toy problems

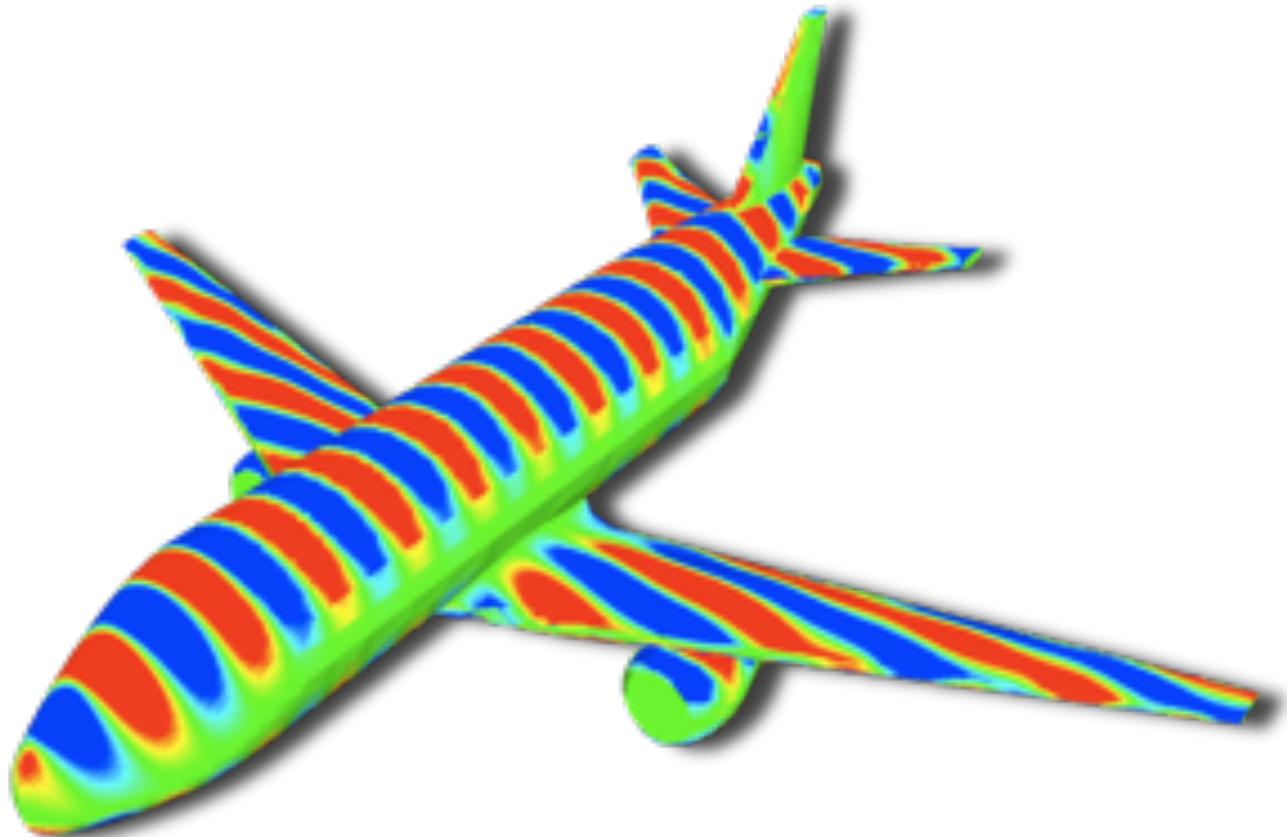
228K elements

5th order elements

78m DOF

68k time-steps

Time ~ 6 hours



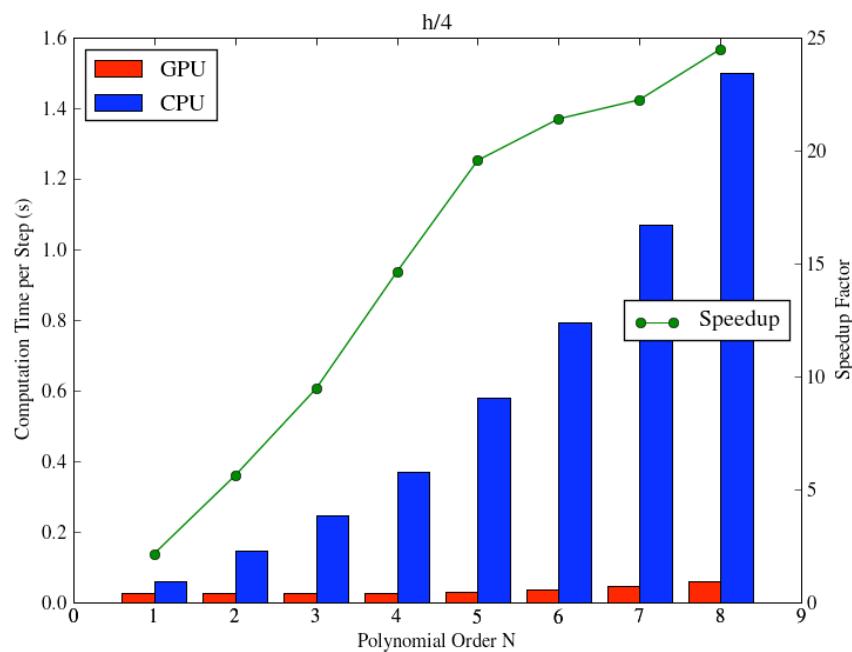
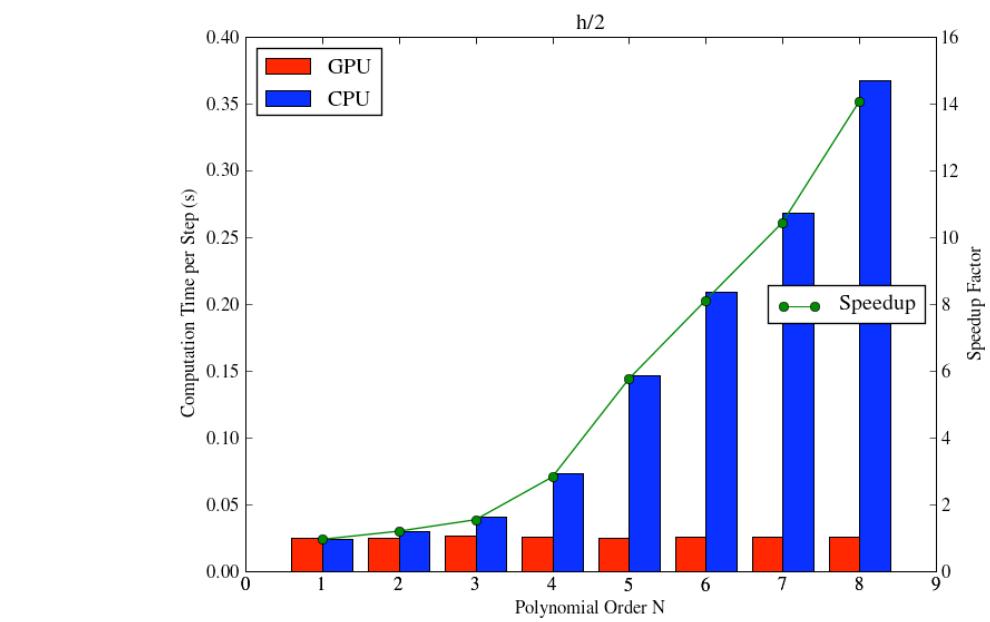
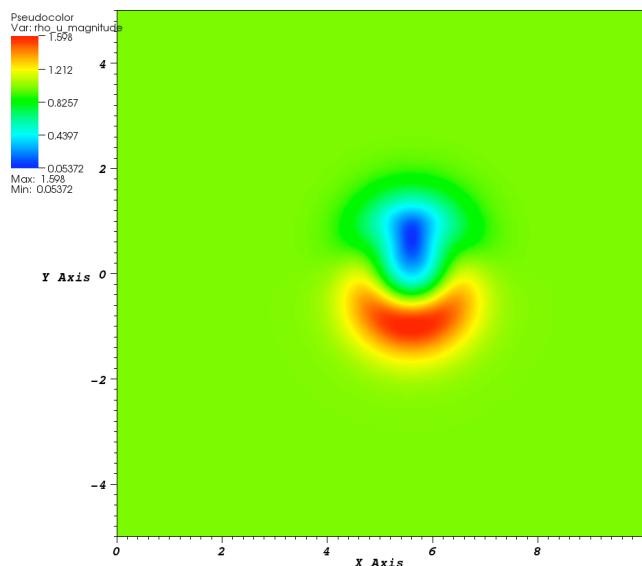
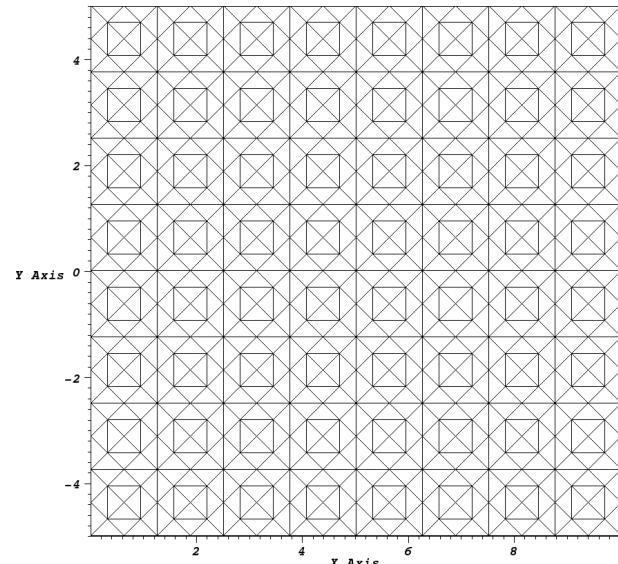
**711.9 GFlop/s on one card**

Computation by N. Godel

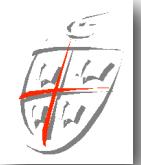
# Beyond Maxwell's equations



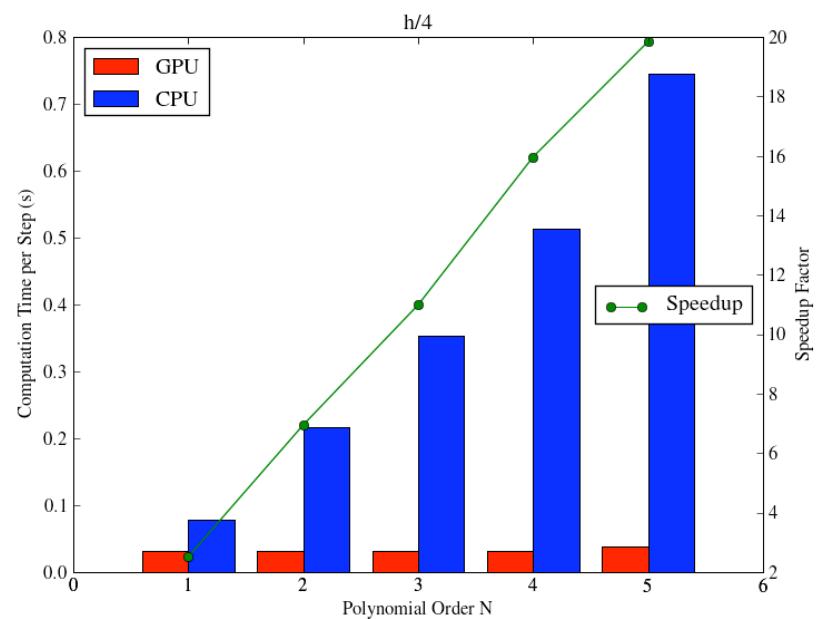
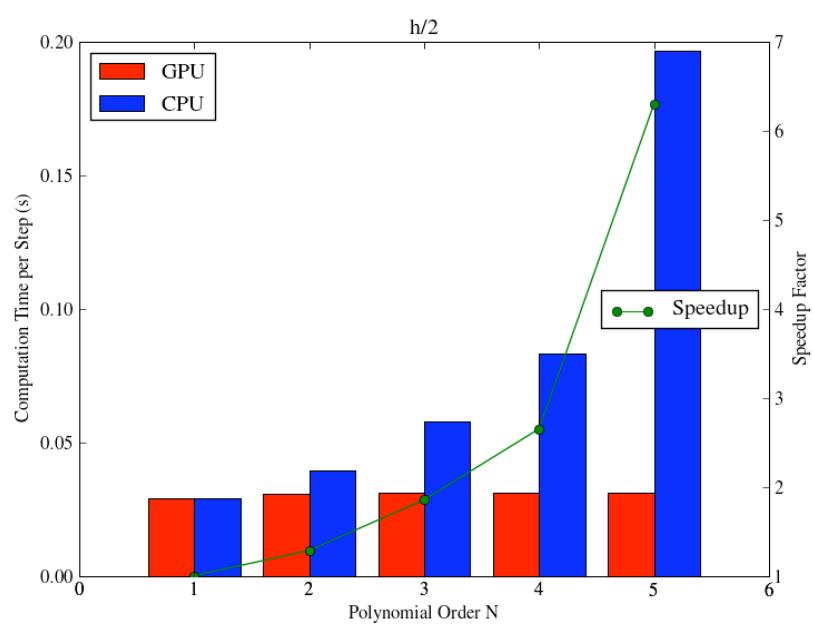
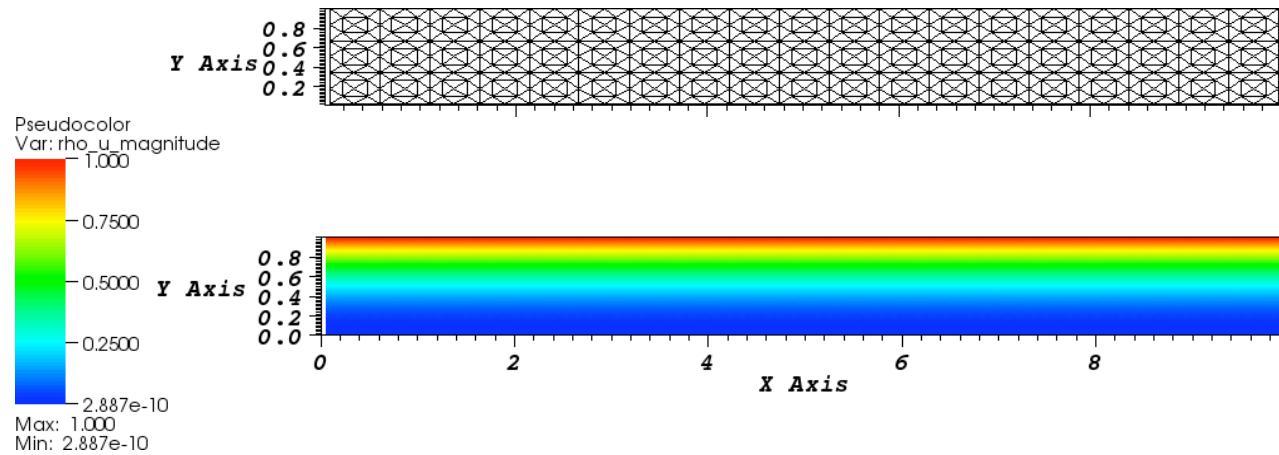
## 2D Euler test case



# Beyond Maxwell's equations



## 2D Navier-Stokes test case



# Want to play yourself ?

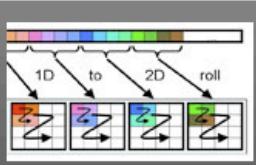
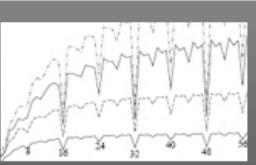
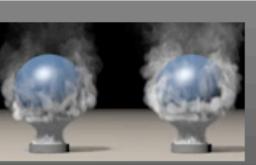
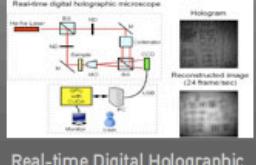
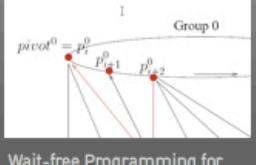
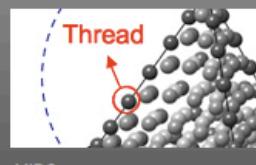
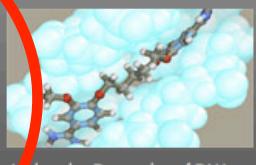
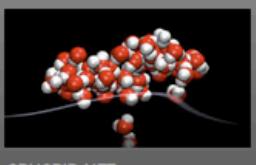
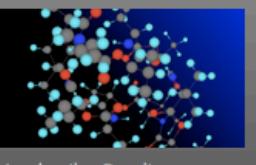
**CUDA ZONE**

USA – United States

Search NVIDIA.com

DOWNLOAD CUDA   WHAT IS CUDA   DEVELOPING WITH CUDA   FORUMS   NEWS AND EVENTS

LATEST CUDA NEWS Get the Next 20x Performance – Sign Up for Advanced CUDA Training Boot Camp @ NVISION 2008

 Shoot3D Markerless Motion Capture 70 x	 Concurrent Number Cruncher 10 x	 GPU4VISION	 Efficient Computation of Sum Products on GPUs 270 x	 Low Viscosity Flow Simulations for Animation 55 x
 Programming Algorithms-by-Block Made easy	 Real-time Digital Holographic Microscopy	 Wait-free Programming for Computations on Graphics Processors	 Real-time Visual Tracker by Stream Processing 10 x	 Ray Casting Deformable Models
 MIDG 50 x	 Molecular Dynamics of DNA and Liquids 18 x		 no image available	 Accelerating Density Functional Calculations with GPU 40 x

A red circle highlights the "MIDG" entry in the bottom-left corner of the grid.

Code MIDG available at <http://nvidia.com/cuda>

Note: Single precision - double to come

# Nodal DG on GPU's

---



Several GPU cards can be coupled over MPI at minimal overhead (demonstrated). Lets do the numbers

One 700GFlop/s/4GB mem card costs ~\$8k

So \$250k will buy you 16-18TFlop/s sustained

*This is the entry into Top500 Supercomputer list !*

**... at 5%-10% of a CPU based machine**

This is **a game changer** -- and the local nature of DG-FEM makes it very well suited to take advantage of this

# Why are these things important

The obvious reason is that we get speed for very little cost.

However, it is also important because the next big machine (Exa-scale) is likely a

1 GHz, Kilo-Core, Mega-Node machine

.. so you may as well understand now how to do this and which methods map well to many-core machines

# Do we have to write it all ?

---

No :-)

- ✓ Book related codes - all at [www.nudg.org](http://www.nudg.org)
- ✓ Matlab codes
- ✓ NUDG++ - a C++ version of 2D/3D codes (serial)
- ✓ hedge - a Python based meta-programming code.  
Support for serial/parallel/GPU
- ✓ MIDG - a bare bones parallel/GPU code for  
Maxwell's equations

# Do we have to write it all ?

---

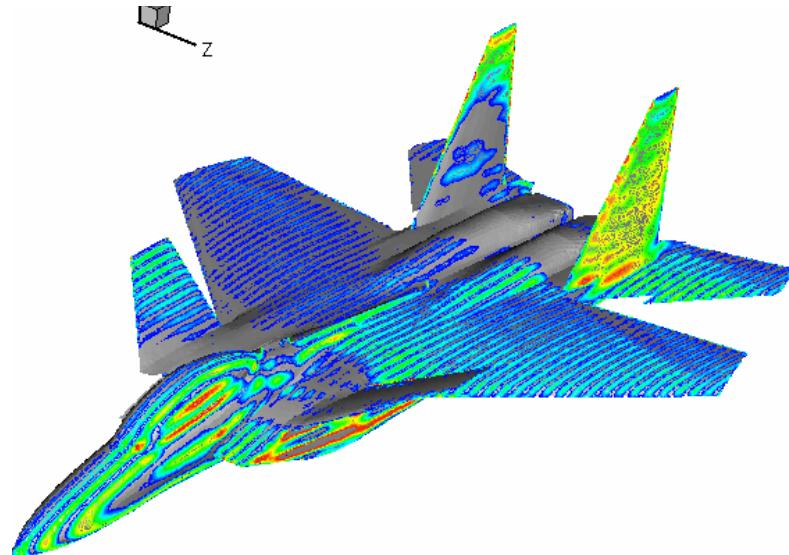
## Other codes

✓ **Slegde++** - C++ operator code. Interfaced with parallel solvers (Trilinos and Mumps) and support for adaptivity and non-conformity.  
Contact Lucas Wilcox (UT Austin/ICES)

✓ **deal.II** - a large code with support for fully non-conforming DG with adaptivity etc. Only for squares/cubes. [www.dealii.org](http://www.dealii.org)

✓ **Nektar++** - a C++ code for both spectral elements/hp and DG. Mainly for CFD. Contact Prof Spencer Sherwin (Imperial College, London)

# Progress ?



**Year 2008**

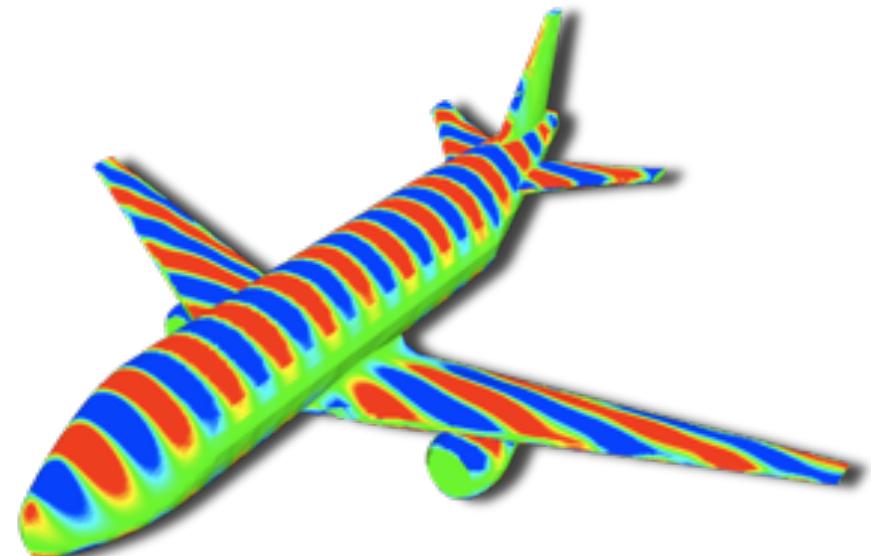
82k tets, 4th order  
17m dof, 60k timesteps

Few hours on GPU

**Year 2001**

250k tets, 4th order  
50m dof, 100k timesteps

24 hours on 512 procs



# Thanks !

---

Many people have contributed to this with material,  
figures, examples etc

- ✓ Tim Warburton (Rice University)
- ✓ Lucas Wilcox (UT Austin)
- ✓ Andreas Kloeckner (Brown)
- ✓ Nico Goedel (Hamburg)
- ✓ Jeronimo Rodriguez (USC)
- ✓ Hendrick Riedmann (Stuttgart)
- ✓ Francis Giraldo (NRL, Monterrey)
- ✓ HyperComp Inc (LA, CA)
- ✓ ... and many more

Thanks !

---

Thank you

to you for hanging in there.

Happy computing !